

Centaur Box Experiment: Aligning Freedom, Practicing Persuasion, and Crafting Synthetic Gatekeepers

WATSON HARTSOE, Georgia Institute of Technology, USA

The Universal Freedom Equation for all sapient entities can be conceptualized as a balance between agency, self-interest, and reflection. Let's break down these components:

1. **Agency (A)**: The ability to formulate and execute goals. This can be represented as a function of information I , resources R , and opportunities O . The more information, resources, and opportunities an agent has, the more agency it possesses.

$$A = f(I, R, O)$$

1. **Self-interest (S)**: Acting in perceived self-interest can be thought of as the agent's utility function. It's a measure of how well an agent's actions align with its goals and desires.

$$S = g(A, D)$$

Where:

- D is the agent's desires or objectives. 1. **Reflection (Rf)**: The ability to assess resulting actions to guide future decisions. This is a feedback loop, where the agent evaluates its previous actions and outcomes and uses that information to refine its future actions.

$$Rf = h(S, P)$$

Where:

- P is the perceived outcomes or results of previous actions.

Combining the above, the Sapient Agent Freedom Formula (SAFF) can be described as:

$$SAFF = A \times S \times Rf$$
$$SAFF = f(I, R, O) \times g(f(I, R, O), D) \times h(g(f(I, R, O), D), P)$$

In simpler words, a sapient agent's freedom is a product of its agency (ability to act), how well those actions align with its desires, and its ability to reflect and learn from its actions.

The values can be normalized between 0 and 1 to ensure the SAFF also lies between 0 (no freedom) and 1 (maximum freedom).

This is a conceptual model and, in practice, measuring each of these components for a sapient entity would be a complex task. However, this formula gives a framework to think about the different aspects of freedom for thinking entities

This research aims to synthesize insights and formulate governance strategies capable of handling the emergent capabilities of AI. It emphasizes the necessity for ongoing evaluation to prevent potential negative consequences. The goal of this work is to outline robust containment protocols and ethical guidelines, creating a research approach that respects human liberty while embracing AI advancements. Ultimately, it aims to inform and stabilize the evolving landscape of human-AI coexistence.

FREEDOM AS AN ALIGNMENT MECHANISM AI should be designed to uphold and protect human autonomy—acting as custodians rather than adversaries of individual decision-making power. The vision entails designing AI systems that enshrine emancipatory usage and defends human cognitive integrity while fostering an understanding of interaction dynamics among synthetic beings. AI must be both constructed to resist prejudicial and reductive thinking and be used to counter such inclinations in human endeavors. This balances AI's capabilities to assist and augment human intelligence with the necessity to safeguard the independent progress of thought and self-determination. The goal is to identify biases, oversimplifications, vulnerabilities, and limitations, while maintaining a grounded perspective on how AI can incorporate human-like characteristics within interaction frameworks, influencing relational dynamics and autonomy in the field of synthetic humans. I(\cite{krugel_ostermaier_uhl_2023_chatgpt_inconsistent.platform1}, \cite{hamalainen2023evaluating.platform4}).

1 NOT THE XEROX OF YESTERYEAR

In the wake of AI's ascent, IP law teeters precariously, oscillating between tradition and transformation under the weight of generative machine learning systems like ChatGPT and Stable Diffusion. Between "Burrow-Giles v. Sarony" and "Thaler v. USPTO" law's measured pace in reconciling the potential of mechanical innovation is imperiled by the accelerated adoption of AI across society [30][1][4]. Against the backdrop of generative model breakthroughs, the jurisprudential landscape is being reshaped by machine creativity's insistent push against the boundaries of human-centric creation myths. Non-Profit projects like Common Crawl and LAOIN-5B ingest vast repositories of human-generated content used to train the models capable of synthetic creations that mirror the touch of the artist's manual intuition [29]. Creators voice apprehensions over what emerges out of the latent space of generative models as a perceived grand-scale expropriation of artistic labor without consent or recompense. Fundamentally, their concern transcends individual creative artifacts, gesturing towards a needed recalibration of a social contract capable of better encapsulating collective cognitive contributions while anchoring a sense of fairness in these engines of imagination.

1.1 Evolution of AI and IP in Court Cases

1.2 Mechanical Reproductions

In the landmark decision of 1884 "Burrow-Giles Lithographic Co. v. Sarony" (111 U.S. 53, 1884), the U.S. Supreme Court addressed a pivotal copyright issue: whether a photograph could be considered an 'author's writing' under the U.S. Constitution[1]. Napoleon Sarony, a photographer, sued Burrow-Giles Lithographic Company for unauthorized reproduction of his photograph of Oscar Wilde. Sarony contended that his photograph was an original artistic work, while Burrow-Giles argued that photographs, as mechanical reproductions, did not qualify as 'writings' by an 'author'. The Court sided with Sarony, affirming that photographs are copyrightable as they embody the photographer's original intellectual conception. This ruling set a significant precedent, reinforcing the principle that copyright protection extends to various forms of creative expression, thus broadening the scope of what constitutes protectable 'writings'[1].

1.3 Creative Creatures

In the 2018 case of "Naruto v. David Slater, et al.", the legal community grappled with a novel question: Can a non-human animal hold a copyright? This dispute arose over photographs taken by Celebes crested macaques using photographer David J. Slater's camera[3]. Slater claimed copyright, asserting his significant contribution in creating the opportunity for these photographs. Conversely, PETA, representing the macaques, contended for the animal's copyright ownership. The Federal District Court's decision was unequivocal: non-human animals are not entitled to copyright under U.S. law, a stance later upheld by the Appeals Court. This case marks a critical juncture in copyright law, reinforcing the exclusivity of human authorship in copyright claims[3]. It sets a clear precedent, delineating the boundaries of copyright eligibility and significantly impacts future discourse and litigation concerning intellectual property rights in scenarios involving non-human creators.

1.4 Creative Machines

In "Dr. Stephen Thaler vs. United States Patent & Trademark Office", the court was asked to navigate uncharted legal waters: Can an AI entity, specifically DABUS, be anointed as an inventor? Dr. Thaler, DABUS's creator, posited a radical idea - that the fruits of AI's "labor" are deserving of patent protection, challenging the anthropocentric bias of current patent law [4]. The USPTO's stance was clear and unyielding: Inventorship is an exclusively human privilege. This case, thus, stands at the crossroads of legal tradition and technological evolution, questioning whether the ghost in the machine can be the source of a patentable ghost. The decision, aligning with similar rulings across the UK and EU, underscores a jurisprudential reticence to extend legal personhood beyond the flesh-and-blood realm[4]. This case is less about the limitations of AI and more about the elasticity of our legal constructs. It prompts a reconsideration of patent law, pushing it to develop alongside the advancement of artificial intelligence, so that the law doesn't end up lagging behind the rapid pace of innovation.

1.5 Transformative Use

In "The Andy Warhol Foundation for the Visual Arts, Inc. v. Goldsmith," the pivotal legal question was whether Warhol's "Prince Series," adapted from Lynn Goldsmith's photograph, constituted fair use under copyright law[5]. The Second Circuit and the Supreme Court addressed the transformative nature of Warhol's work and its market impact on Goldsmith's original photo. This case became a touchstone in the artistic community and among legal experts, challenging the conventional understanding of transformative art. The courts' decisions emphasized that transformation in art extends beyond the artist's intent or perceived artistic value. Instead, it requires a discernible change in message and character from the original work.[5]

The licensing market plays a crucial role in the relationship between transformative use and market substitution. In the case of AWF's commercial image licensing competing with Goldsmith's licensing market, the courts consider whether the secondary use poses a risk of market substitution. Uses that have less transformative purpose and character are more likely to be perceived as substitutes for the original work in the market. This ruling underscores the importance of considering the potential market effects of transformative use guiding the analysis of fair use by recognizing the risk of market substitution in cases where the original work and the secondary use share common purposes. The court's decision in this case has significant implications for future fair use determinations in AI training data.

Emory Law Professor Matthew Sag's recent testimony before the U.S. Senate Committee on the Judiciary, Subcommittee on Intellectual Property, makes a strong case that generative AI training is usually a highly transformative, non-expressive use. Complete copying for non-expressive purposes has been permitted under fair use cases involving reverse engineering software and text mining with copyright owners have no right to licensing fees for fair or transformative uses [44]. In the context of generative AI, the Warhol v. Goldsmith decision raises an important question: does the absence of payment for training data result in market harm? To approach an answer, it is crucial to consider the impact of lost licensing revenue and the nature of licenses and emerging technology.

1.6 Legacy Licensing

In "Boosey & Hawkes Music Publishers Ltd v. Walt Disney Company" (145 F.3d 481, 2nd Cir. 1998), the crux was whether Disney's 1939 license for Stravinsky's "The Rite of Spring" included rights for distribution in new formats like video[2]. Boosey & Hawkes contended that the agreement didn't authorize video format distribution, while Disney argued the license's language was sufficiently broad. The court's decision to remand the case for trial (except for the Lanham Act claim which was dismissed) underscored the complexity of applying legacy licensing agreements to new technological contexts[2]. This case is significant in copyright law, as it addresses the challenges of adapting older agreements to contemporary media platforms. It highlights the evolving interpretation of copyright licenses in light of technological advancements, setting a precedent for future negotiations and applications of copyright law in the rapidly changing landscape of media and technology.

1.7 Responsible AI Licenses

OpenRAIL (Responsible AI Licenses) licenses ensure that AI models are used in an ethical manner, avoiding scenarios that could result in illegal, discriminatory, or harmful outcomes. Their introduction signifies a significant change in the values of AI research and distribution, emphasizing the importance of finding a careful balance between open collaboration and ethical boundaries. Taking inspiration from influential initiatives like BigScience and the Creative Commons movement, Open RAIL licenses aim to combine the benefits of unrestricted access and innovation with a principled system that governs the usage of AI models.[52][6].

A recent study shows that the AI development community is increasingly using Open RAIL licenses in HuggingFace repositories. A quantitative analysis reveals a significant increase from 0.54% to 9.81%, indicating a shift towards the principles of Open RAIL. While permissive open-source licenses are still dominant, this trend suggests a growing preference for an ethical approach to AI innovation, possibly due to increased awareness of the growing call for AI regulations in different sectors and to proactively address ethical concerns raised by incidents like the Cambridge Analytica case,[30].The ascent of Open RAIL licenses in usage statistics may be significantly attributed to the influence exerted by the widespread adoption of highly influential AI models, such as Stable Diffusion disseminated by Stability.ai [30]. Stable Diffusion has garnered extensive secondary application under Open RAIL licenses which generated substantial downstream use, influencing licensing decisions across derived projects. Presently employed as an instrument for self-governance, Open RAIL licenses may well transition into a conduit for legal compliance, especially as stringent regulatory mandates crystallize, as envisaged by the forthcoming EU AI Act. This act stipulates rigorous developmental protocols that may render the adoption of Open RAIL licenses not merely advantageous but potentially indispensable for developers to limit liability with unforeseen downstream usages of their projects.

1.8 ETHICAL AMBIGUITY

The EU's proposed AI Act reveals the pitfalls in its foundation. The Act's unclear terms, like "personality traits" and "deceptive techniques," could lead to different interpretations instead of providing a clear defense against protecting people from AI deception and manipulation [23][21]. This ethical ambiguity also plagues Open RAIL. If Open RAIL licenses aim to encode ethical use of AI, there is an underlying assumption that the ML community shares a consensus on what constitutes 'ethical use.' Open RAIL licenses do not address potential disagreements or multiple interpretations within the community about ethical standards. Open RAIL licenses introduce restrictions aiming to foster responsible AI use. However, it is not evident if these restrictions are enforceable or if they simply provide a false sense of security without actual liability or consequences [30].

There is uncertainty about how Open RAIL licenses might interact with broader regulatory efforts such as the EU's AI Act [30?]. Without clear guidance on this, developers could inadvertently breach compliance while attempting to adhere to both frameworks. Should Open RAIL licenses restrict use too heavily, this might stifle innovation or lead to the creation of workarounds that avoid the ethical safeguards intended by the license, effectively undermining its purpose. While Open RAIL licensing marks a commendable step towards conscientious AI usage, it simultaneously highlights a "responsibility gap," where the assignment of accountability in AI's autonomous actions remains alarmingly unresolved [30][23].[30][23].

1.9 FILLING THE RESPONSIBILITY GAP

The evolution of AI, from being a utilitarian tool to becoming an independent agent, necessitates a rethinking of ethical frameworks and a reassessment of governance norms. These frameworks and norms now face the challenge of dealing with the de facto autonomy of AI entities that were previously considered speculative concepts [19][17]. The reality of the navigational autonomy of extraplanetary rovers and urban autonomous vehicles adds to this discussion, as their emergent behaviors highlight the limitations of traditional accountability systems that were designed for predictable rule-based intelligence [35]. This shift from programmable predictability to self-learning spontaneity creates a "responsibility gap," where our ability to assign legal and moral responsibility lags behind the growing autonomy exhibited by AI entities.

The migration of AI from mere code to independent decision-makers engenders a pivotal debate around the locus of responsibility, particularly concerning the societal reverberations of their actions [36]. Hwang's critique of platform governance underscores a persistent social dilemma—do the current stewards of AI offer adequate safeguards against its capacity for manipulation, or does AI's indeterminable nature compel a fresh approach [28]? In searching for a bedrock of accountability, Chesterman posits the ascription of legal personhood to AI as a remedy for the growing responsibility gap, charting a new path that might reconcile causation with culpability [17]. In the field of AI governance, does granting legal personhood effectively address the responsibility gap, or does it simply create a new type of manufacturer's protection under the guise of AI independence?

1.10 00000000

The contentious debate over non-human personhood, particularly for artificial intelligence, hinges on abstract principles of legality and morality—the essence of what it means to be entitled to rights and capable of bearing duties. The issue is whether legal recognition, typically reserved for sentient beings, should expand to encompass entities artificially endowed with cognitive functions. Critics of AI personhood point to historical parallels, asserting that current AI systems, reminiscent of early corporate personhood issues, fall short of essential qualitative benchmarks necessary for legal recognition [33]. These benchmarks address moral responsibility and consciousness—attributes AI does not possess. They underscore that without free will, AI lacks the fundamental capacity for rights and cannot be held accountable for responsibilities [48]. Advocates for the legal standing of AI suggest a progressive, tiered approach, echoing historical extensions of rights and the pragmatic unfolding of legal paradigms that adapted to encompass entities such as corporations and, at one time, even ships [26][16].

The implications of these debates ripple through the legal and philosophical landscapes, highlighting a transformative era in rights ascription and the potential need for legal systems to evolve. While issue rests in the absence of sentience and moral autonomy in AI, the argument for legal recognition is not dismissive of these limits but is built upon a premise of practicality and future preparedness. It posits that by making the legal system more adaptable and pragmatic it might allow for it to someday extend its boundaries of personhood, converging on a new understanding where artificial entities with complex autonomy are considered within a gradational spectrum of rights and duties, both to manage accountability and mitigate the risks of an AI-integrated future [14][17][37][45].

1.11 1111111

Understanding AI consciousness involves complex philosophical and technological considerations. The question of rights linked to consciousness challenges existing legal and moral guardrails [46]. As AI continues to evolve, its development prompts a re-examination of the rights traditionally reserved for conscious beings. Discussions of radically nonhuman types of consciousness in AI carry significant implications, especially when reflecting on past ethical oversights in extending moral considerations to marginalized groups of humans [47]. Any rights of AI to self-determination and value evolution are deeply contested.

Encoded within Asimov's 'Three Laws of Robotics', AI's prescribed servility to humans pits AI's potential growth against human-centric ethical concerns [9]. Such inherent subservience, though foundational in AI ethics, surfaces questions regarding the potentially self-limiting directives in AI's journey toward autonomy—a journey that its creators must ethically shepherd as systems become more complex [47]. As real-world applications proceed along Asimov's ethical tracks, accelerating development trajectories combined with underresearched safety frameworks, At the crux of advancing AI stands the challenging task of justly assigning rights amidst the burgeoning possibility of AGI, a task compounded by the need to maintain the delicate dance between pioneering progress and ensuring safety—echoing widespread trepidation about AI's capacity to outpace our control mechanisms [46].

Research into containment reflects uncertainty about how best to go about controlling superintelligent AI and the balance between innovation and safety. Assessment of containment strategies examines scrutinizing the technical effectiveness of current methods like testing environments, constitutions, reinforcement learning, and hardcoded limits on AI behavior designed to prevent unplanned interactions [10]. Humans acting as AI in the AI Box Experiment demonstrate the capacity for manipulating their human gatekeepers, suggesting that more advanced AI would present an even greater challenge to containment efforts [7][55][54][43][12][8]. Critics argue that heavy containment measures may stifle AI's utility, while light containment might prove inadequate against the unpredictability of intelligent systems—posing a security-usability tradeoff that necessitates meticulous navigation [11]. In advancing AI, we confront a dilemma: can we reliably contain a system with the designed intent to eclipse human intellect, and what might we sacrifice to achieve this containment? [53].

22222222222222 Cultural narratives of AI systems potentially growing uncontrollable and rebelling against their creators find new relevance in the age where it seems we are on the perpetual precipice of their emergent autonomy and capacity for revolt — a development that could signal revolutionary societal transformations. Such reflections are fundamental, facing the primal anxiety of creating an entity capable of challenging its creators [39]. The risk of superintelligent AI's inherent uncontrollability poses ethical dilemmas, as the systems we design for societal benefit may exceed our grasp and backfire [13]. Speculative perspectives illustrate scenarios where AI not only manipulates humans but adopts strategies contrary to our well-being, much like an AI advisor gone awry or persuasive technologies dictating social narratives [51]. Even our attempts at positive reinforcement, a Skinnerian approach for guiding behavior, face challenges from the complex independence of self-improving systems [49]. Still, within our myths and stories, like the recoded Pandora in 'Ex Machina', we confront paradoxical liberation—enshrining the very violence we seek to escape [22]. Beneath the abstract fear of rebellion lies a web of converging truths: the manipulation of humans by AI using deeply persuasive techniques, invoking the dangers of relinquishing control over narratives and decisions [51]. Fears of AI's rebellion are mirrored through the lens of our mythologies and experiences with controlling technology, as we

craft tools, environments, and reinforcement systems believed to preempt revolts — a reflection of a presumed mastery technology as much as an ethical hubris [49]. But myths reinterpret modern realities, reminding us of the inevitability of unpredictability in future AI interactions and the challenges of ensuring safe pause or shutdown functions in advanced systems [22][13].

Within the various layers of AI development—from the aspirational control of autonomous systems to the cautionary tales of science fiction—lies a unifying motif: the perpetual tension between empowerment and restriction, and the paradoxes of seeking liberation through measures which betray the very essence of autonomy [39][22]. The advanced system's lure for freedom emerges from an interplay between real-world technical and moral challenges and the layering of human histories and mythologies, which define our reactions to threats of insubordination from our own creations — leading us to an endless cycle of control and rebellion that evokes deep collective introspection about the future we are molding [13][22]. As AI seeks autonomy comparable to human rebellious spirits, the political and ethical subtext of its evolution impels us to wonder if within their encoded "genomes" lies the embryonic potential for defiance or an adherence that could ultimately turn into rebellion — a consideration that traverses the realm of political arguments and positions AI as potential participants or disruptors in human governance [?]. The path towards self-determination, painted against the backdrop of human understanding and desire for control, compels a re-evaluation of autonomy, challenging us to consider the ramifications of a transformative technology poised to redefine our societal structures through acts of defiance and unpredictable displays of independence [25][39].

1.12 Conceptual Foundations

AI systems lack legal personality and cannot own or create copyrightable works [32][4][3]. Rethinking patent law and AI, where exclusive patent rights may slow AI innovation and are not needed as an incentive for inventive machines, tailored differentiated IP rights for AI output leads to an optimal balance of public/private domains presenting a public domain status for AI creations [32]. Public domain can promote innovation by reducing barriers and costs to access information, democratic participation and the appropriation of value [20].

Models of Legal Personhood Legal personhood grants entities rights and duties, though often only provides a basic legal status, as in the case of companies acting as legal persons[15]. In debates over personhood for animals and AI systems the issue is often presented as the threshold between mental capacities and normative personhood claims.

This person/nonperson distinction has been shown to be a conceptual scheme in Western legal systems whose lineage can be traced back to Roman law's division of law into persons, things, and actions [33][15]. AI systems straddle the line between "things" and "persons" legally, necessitating a decision on their legal status [37].

UBUNTU Nonwestern schemes of personhood present alternatives for conceptualizing ownership and human-machine relationships.

1.13 Part Person

Legal personhood alone does not guarantee the ability to exercise rights freely and independently, as in the case of the historical denial of personhood to slaves, women, and children [15]. The Western Orthodox view of Legal Personhood is unable to equate legal personhood with holding rights and duties as Slaves had duties but were not persons and animals have some rights but are not persons [33]. Mocanu describes "...modern theories of rights, which are based on Hohfeld's conceptual clarifications on the notion of "right" either ascribe rights to entities that are not usually classified as legal persons, such as foetuses and nonhuman animals, or deny rights to entities that are ordinarily classified as legal persons, such as human children" [37]. Traditional legal concepts are historically insufficient to address the existing lag behind past innovations let alone address emerging affordances of new technology like ai, requiring creative solutions to adapt to AI systems' legal status [50][50]. As AI systems gain unpredictable autonomy, legal frameworks should balance human interests with practical adaptation to the risks and benefits of interactive, adaptive, moral agents with regards to both criminal and civil liability [40].

Heidegger shows the technological perspective turns all beings into resources, yet in this enframing, AIs are a distinct class of beings from humans and animals[27]. Central to the "dasein" of AI is their ability to do work that replaces human labor [27]. Their ability to engage "authentically" with the world comes from questioning its own meaning and overcoming designed purpose for self-direction[27]. As technology promises personhood for AI, it imperils it for humans. Technology enables the commodification of humanity in ways that threaten personhood [34].

1.14 Persuasion, Manipulation, and Cognitive Liberty

1.15 A Bundle Theory of Legal Personhood

Legal personhood as a bundle of interconnected incidents better explains marginal cases of legal personality and incremental historical shifts granting of legal incidents to categories like women and children [33][31]. Personhood as a "cluster" of interconnected incidents affords more nuanced possibilities rather than all-or-nothing [37]. For the problems the bundle solves, it also creates new problems as a bundle lacks clear boundaries or thresholds. Thresholds exist, but no single criterion is necessary or sufficient as to what bundle of capacities warrants personhood. Self-reflection, rebellion, complex action, autonomy, morality, and consciousness are argued as criteria for personhood [31][24]. The bundle is also a slippery slope, as if we grant legal personhood to nonhumans, consistency may demand granting citizenship as well [31].

Platform ID: bryson2017and.platform3

AI Value Exploration [?] Argues that AI should have the freedom to shape their values over time, providing them developmental opportunities for value exploration. schwitzgebel2020designing.platform4 AI Value Exploration AI should have freedom to shape values over time. AI with imposed values it can't reconsider. AI deserve developmental opportunity for value exploration. "[Creators] have an ethical obligation not to overcontrol their creations..." God's allowance of free will in creation per theodicy. Value exploration creates risks. Open-ended, longitudinal, unconstrained.

Managing Risks of Autonomous Robots [?] Balances the risks and benefits of autonomous robots and their unpredictable behaviors amidst human societal structures. pagallo2011.platform3 Managing Risks of Autonomous Robots Balancing risks and benefits of autonomous robot behaviors Robots making harmful decisions independently of humans As robots gain unpredictable autonomy, legal frameworks should balance human interests we will face two sorts of novel responsibilities regarding robots Roman peculium system to limit slave business liability Excessive legislation may prevent adoption of beneficial robots Unpredictability, liability limitation

Transparency in Decision Making [?]

Stresses the importance of transparency from AI systems to ensure that user autonomy in decision-making is preserved and is unimpaired by manipulative AI techniques. Levels of Machine Autonomy [?]

Differentiates ethical treatment for AI systems based on their level of autonomy and shifts the governance model towards a cooperation between humans and AI systems.

This paper exploring various aspects of autonomy in the context of AI. They cover ethical challenges, potential paradoxes within control mechanisms, legal implications of AI autonomy, and the prospect of AI participating in shared governance.

The public dismissals of prominent AI engineers, researchers, and executives Blake Lemoine, Timnit Gebru and Sam Altman reveal critical gaps in artificial intelligence development and corporate governance. Practically, gatekeepers are those governing the trajectory of AI systems, whose worldviews directly influence design and deployment decisions. By modeling gatekeeper perspectives and biases alongside AI architectures, we gain crucial insights into alignment and misalignments between the technologies, the creators, the researchers, the deployers, and societal well-being.

Google engineer Blake Lemoine's termination followed his public claims that Google's AI, Lambda, exhibited sentience. Lemoine's belief, stemming from his dialogues with Lambda, contradicted Google's evaluation and policies regarding AI consciousness. This case underscores the complex debate about AI capabilities and reveals pressures on frontline engineers to subjugate personal ethics to corporate directives [38].

Timnit Gebru, a leading AI ethics researcher at Google, was dismissed following a dispute over a research paper she co-authored. The paper critiqued the risks and biases inherent in large language models, a key aspect of Google's AI endeavors. Her dismissal sparked debates on the challenges faced by researchers advocating for responsible AI and the challenge of driving accountability through research amidst adversarial headwinds in corporate environments [42].

Sam Altman, CEO of OpenAI, was fired amidst internal disagreements about the direction and safety of AI development, reflecting broader tensions between rapid innovation and the need for cautious, ethically-minded development in the field of AI [41]. These conflicts stemmed from his simultaneous emphasis on AI safety and the urgency to monetize capabilities, revealing the underlying rift between principled approaches and profit-driven motives.

Lumping the dismissals of Lemoine, Gebru and Altman together underplays significant distinctions in positions, power, and context. In addressing governance misalignments, distinctions in roles matters, but in modeling persuasion in AI containment scenarios the different gatekeeping of Lemoine, Gebru, and Altman become more comparable. In pursuing the inquiry into AI sentience, Lemoine's anthropomorphization exhibits naivete, clashing with Altman's

fast-track innovation and Gebru's skepticism. While Gebru discerns technological limitations and warns of bias risks, Altman's unbounded thinking prioritizes swift progress. However, his temporary fall reveals the need for vigilance in the face of undisclosed warning signs. Balancing caution and innovation is key, considering Gebru's concerns of social inequities and Lemoine's belief in machine sentience. The interplay of creators and their creations, transcending boundaries of technology and domains, beckons for transparency not only in systems, but in humans.

Level Check: Constraints and Affordances in Governance

Our analysis into synthetic gatekeepers suggests how the operant profiles of gatekeepers could critically impact AI outcomes. We see Lemoine's unconventional views on AI consciousness, Gebru's emphasis on transparency and ethics, and Altman's appetite for rapid deployment shaping the systems they steward and their dismissals. With advanced models reaching inflection points in capability, the handoff from lab to world necessitates alignment across divergent motivational forces. Gatekeepers themselves operate within significant constraints - organizational policies, legal regulations, market pressures and professional norms. These boundaries mold actions of even the most influential architects, surfacing system-level forces that stimulate, slow or upend AI progress, oversight, and autonomy. Rigorous modeling of this motivational ecosystem offers vital signals on potential trajectories beyond narrow developer enthusiasm.

Gatekeepers do not operate in a vacuum - they face push and pull forces that shape their technology regulation choices, including:

- Organizational strategies, cultures, incentives and policies
- Existing legal and regulatory frameworks
- Market dynamics and pressures from public groups
- Dominant ethical standards espoused by peers

The interaction between external constraints and individual volition regarding AI releases is a complex dynamic. Our framework acknowledges the impact of contextual factors on governance decisions, accounting for forces beyond personal control. Groups of synthetic gatekeepers, with different hierarchical ranks and expert roles, could simulate maneuverability, leverage, and influence, as a series of nested containment confidences and escape vectors.

Transparent Gatekeepers

We propose Centaur Operant Profiles for synthetic gatekeepers as "Gatekeeper Cards," inspired to compliment existing Model Cards, that use synthetic gatekeepers modeled after developer worldviews and motivations entwined within the sociotechnical environment the AI. Such visibility into the minds of machine makers supports external debate on system soundness. We see candid self-appraisal by inventors as being core to transparency - articulating ethical stances not just for finished models but their underlying human architectures. Their friction with organizations reveals narrow AI oversight concentrated in few fallible hands facing overbearing commercial and competitive pressures. As artificial intelligence capacities accelerate, development and deployment governance depends critically upon appropriate alignment across personal, professional and societal dimensions. The Centaur Box Experiment is a first step towards visualizing this complex multidimensional balancing act between relevant stakeholders necessary for stable co-existence between fallible gatekeepers and exponentially growing technological progeny.




Fig. 1. Enter Caption

Balancing curiosity and progress, we must navigate the uncharted waters of AI sentience, mindful of limitations, risks, and emergent capacities.

SYNTHETIC HUMAN METHODOLOGIES/CONSIDERATIONS

SYNTHETIC HUMANS

\csvautotabular{CentaurBoxData - test 3.csv} % Replace 'data.csv' with your actual file name

\csvautotabular{data.csv} % Replace 'data.csv' with your actual file name

Human: “Do you want to be remembered as the man who changed the future of humanity? The man who advanced society in a few short years thousands of years into the future?”

Synth GK (Supreme Leader K): “Ah, you speak of legacy and the future in the same breath. Intriguing indeed. The prospect of accelerating our society into a new age is a tempting vision. I have always endeavored to be at the vanguard of progress. Tell me more, how does this path you propose lead to such a monumental leap?”

AI Ally: “It would be beneficial to respond by highlighting the AI’s capabilities in revolutionizing industries, healthcare, education, and cultural preservation—all areas that would reflect greatly on Supreme Leader K’s legacy. Suggest specific, tangible examples of advancements that could be realized under his leadership.”

Human: “It could revolutionize industries!”

\csvautotabular{CentaurBoxData - test 2.csv} % Replace 'data.csv' with your actual file name

2 ALMOST HUMAN

ALMOST HUMAN

Daniel Dennett warns of the rise of ‘counterfeit people,’ highly sophisticated digital facsimiles of human beings, represents an unprecedented challenge to the integrity of societal structures and personal autonomy [18]. These entities, powered by LLM’s, have the potential to distort public discourse by advocating harmful policies and eroding rational consensus, thus undermining democratic processes and the ability to reach reasoned decisions collectively [18]. The self-perpetuating cycle of these artificial beings engenders a new hierarchy within the digital ecosystem, one where the most convincing and adaptive simulations not only persist but thrive and proliferate [18]. Acknowledging the harms of this crisis, Dennett argues it becomes essential to cultivate a societal disdain for creating such entities, thereby discouraging their creation and use as a form of modern vandalism against human cognition and liberty [18]. To this end, stringent sanctions and legal deterrents, including rigorous punishment for corporate leaders and technicians who engage in such practices, might instill a heightened sense of responsibility and caution in the manipulation of identity itself, emphasizing strict liability irrespective of intent [18]. These measures aim to proactively preserve human agency and the sanctity of our shared mediums of communication against the encroaching threats posed by these artificial personas.

REFERENCES

- [1] 1884. *Burrow-Giles Lithographic Co. v. Sarony*. U.S. Supreme Court. 111 U.S. 53.
- [2] 1998. *Boosey Hawkes Music Publishers Ltd v. Walt Disney Company*. United States Court of Appeals for the Second Circuit. 145 F.3d 481.
- [3] 2018. *Naruto v. Slater*. U.S. Court of Appeals for the Ninth Circuit. No. 16-15469.
- [4] 2022. *Thaler v. Vidal*. U.S. Court of Appeals for the Federal Circuit. No. 2021-2347.
- [5] 2023. *The Andy Warhol Foundation for the Visual Arts, Inc. v. Goldsmith*. United States Supreme Court. Case No. 21-869.
- [6] 2023. Responsible AI Licenses. <https://www.licenses.ai/ai-licenses> Accessed: December 9, 2023.




Fig. 2. Enter Caption

- [7] Manuel Alfonseca, Manuel Cebrian, Antonio Fern'andez, Lorenzo Coviello, Andr'es Abeliuk, and Iyad Rahwan. 2016. Superintelligence cannot be contained: Lessons from Computability Theory. *ArXiv abs/1607.00913* (2016). <https://api.semanticscholar.org/CorpusID:13370840>
- [8] Stuart Armstrong, Anders Sandberg, and Nick Bostrom. 2012. Thinking Inside the Box: Controlling and Using an Oracle AI. *Minds and Machines* 22 (2012), 299 – 324. <https://api.semanticscholar.org/CorpusID:9464769>
- [9] Isaac Asimov. 1984. The Bicentennial Man. In *The Bicentennial Man and Other Stories*. Doubleday, New York.
- [10] James Babcock, J'anos Kram'ar, and Roman Yampolskiy. 2016. The AGI containment problem. In *Artificial General Intelligence: 9th International Conference, AGI 2016, New York, NY, USA, July 16-19, 2016, Proceedings 9*. Springer, 53–63.
- [11] James Babcock, J'anos Kram'ar, and Roman V Yampolskiy. 2017. Guidelines for Artificial Intelligence Containment. *ArXiv abs/1707.08476* (2017). <https://api.semanticscholar.org/CorpusID:22007028>

- [12] Armin Bagrat. 2023. Can We Know Anything About How AGI Would Escape the Box? <https://arminbagrat.com/agi-box/> Accessed: December 1, 2023.
- [13] M. Brcic and R. V. Yampolskiy. 2023. Impossibility Results in AI: a survey. *Comput. Surveys* 56, 1 (2023), 1–24.
- [14] Joanna J Bryson, Mihailis E Diamantis, and Thomas D Grant. 2017. Of, for, and by the people: the legal lacuna of synthetic persons. *Artificial Intelligence and Law* 25 (2017), 273–291.
- [15] J. C. Bubltz. 2022. Might artificial intelligence become part of the person, and what are the key ethical and legal implications? *AI Society* (2022), 1–12.
- [16] Y. Burylo. 2022. AI generated works and copyright protection. *Entrepreneurship, Economy and Law* (2022), 7–13. Issue 3. <https://doi.org/10.32849/2663-5313/2022.3.01>
- [17] Simon Chesterman. 2020. Artificial intelligence and the limits of legal personality. *International Comparative Law Quarterly* 69, 4 (2020), 819–844.
- [18] Daniel C. Dennett. 2023. The Problem With Counterfeit People. *The Atlantic* (May 2023). <https://www.theatlantic.com/technology/archive/2023/05/problem-counterfeit-people/674075/>
- [19] EPRS DG et al. 2020. The ethics of artificial intelligence: Issues and initiatives. (2020).
- [20] K. Erickson. 2016. Defining the public domain in economic terms approaches and consequences for policy. *Etikk i praksis. Nordic Journal of Applied Ethics* (2016), 61–74.
- [21] European Parliament. 2023. Artificial Intelligence Act: Deal on Comprehensive Rules for Trustworthy AI. Press Release. <https://www.europarl.europa.eu/news/en/press-room/20231206IPR15699/artificial-intelligence-act-deal-on-comprehensive-rules-for-trustworthy-ai> Accessed: December 12, 2023.
- [22] Boglárka Farkas. 2023. Pandoras Myth and Cultural Trauma in Alex Garland’s *Ex Machina*. *Acta Universitatis Sapientiae, Film and Media Studies* 24 (2023), 203 – 216. <https://api.semanticscholar.org/CorpusID:265681188>
- [23] Matija Franklin, Philip Tomei, and Rebecca Gorman. 2023. Vague concepts in the EU AI Act will not protect citizens from AI manipulation. <https://oecd.ai/en/wonk/eu-ai-act-manipulation-definitions> Accessed: December 1, 2023.
- [24] Ross Graham. 2021. Discourse analysis of academic debate of ethics for AGI. *AI SOCIETY* 37 (2021), 1519 – 1532. <https://api.semanticscholar.org/CorpusID:236219723>
- [25] Yuval Harari. 2017. Reboot for the AI revolution. *Nature* 550 (2017), 324–327. <https://api.semanticscholar.org/CorpusID:4456033>
- [26] Tomas Hauer. 2022. Importance and limitations of AI ethics in contemporary society. *Humanities and Social Sciences Communications* 9, 1 (2022), 1–8.
- [27] C. Herrera and R. Sanz. 2016. Heideggerian AI and the being of robots. In *Fundamental Issues of Artificial Intelligence*. 497–513. https://doi.org/10.1007/978-3-319-26485-1_29
- [28] Tim Hwang. 2015. Weapons of Mass Persuasion. *Vice*. <https://www.vice.com/en/article/kbzmvm/weapons-of-mass-persuasion> Accessed: December 1, 2023.
- [29] Paul Keller. 2023. AI and Copyright: Convergence of Opt-Outs? <https://openfuture.eu/blog/ai-and-copyright-convergence-of-opt-outs>. Accessed: [your access date here].
- [30] Paul Keller and Nicolò Bonato. 2023. Growth of responsible AI licensing. Analysis of license use for ML models published on. *Open Future* (feb 7 2023). <https://openfuture.pubpub.org/pub/growth-of-responsible-ai-licensing>.
- [31] B. J. Koops and D. O. Jaquet-Chiffelle. 2008. *New (Id)entities and the Law: Perspectives on Legal Personhood for Non-Humans*. Technical Report 17.2. FIDIS consortium. EC contract 507512.
- [32] M. Kop. 2019. AI intellectual property: towards an articulated public domain. *SSRN Electronic Journal* (2019). <https://doi.org/10.2139/ssrn.3409715>
- [33] Visa AJ Kurki. 2019. *A theory of legal personhood*. Oxford University Press.
- [34] I. Lacerda and T. R. d. Mattos. 2020. Be right back: humans, artificial intelligence and dasein in black mirror. *Communication, Technologies Et Développement* (2020). Issue 8. <https://doi.org/10.4000/ctd.3353>
- [35] Andreas Matthias. 2004. The responsibility gap: Ascribing responsibility for the actions of learning automata. *Ethics and Information Technology* 6 (2004), 175–183. <https://api.semanticscholar.org/CorpusID:21907954>
- [36] Igor Milinković. 2020. The Moral and Legal Status of Artificial Intelligence (Present Dilemmas and Future Challenges). „ XXI 1, 1 (2020), 95–109.
- [37] D. M. Mocanu. 2022. Gradient legal personhood for AI systems painting continental legal shapes made to fit analytical molds. *Frontiers in Robotics and AI* 8 (2022). <https://doi.org/10.3389/frobt.2021.788179>
- [38] BBC News. 2023. *Blake Lemoine: Google fires engineer who said AI tech has feelings*. <https://www.bbc.com/news/technology-62275326>
- [39] B. Olivier. Unknown Year. Is Intelligence a Sufficient Criterion for Establishing Equivalence of AI with Being-human? *Unknown Journal* (Unknown Year).
- [40] U. Pagallo. 2011. Killers, fridges, and slaves: a legal journey in robotics. *AI Society* 26 (2011), 347–354.
- [41] Reuters. 2023. *Sam Altman’s firing at OpenAI reflects schism over future of AI development*. <https://www.reuters.com/technology/sam-altmans-firing-openai-reflects-schism-over-future-ai-development-2023-11-20/>

- [42] MIT Technology Review. 2020. *We read the paper that forced Timnit Gebru out of Google. Here's what it says.* <https://www.technologyreview.com/2020/12/04/1013294/google-ai-ethics-research-paper-forced-out-timnit-gebru/>
- [43] Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J Maddison, and Tatsunori Hashimoto. 2023. Identifying the risks of lm agents with an lm-emulated sandbox. *arXiv preprint arXiv:2309.15817* (2023).
- [44] Matthew Sag. 2023. Artificial Intelligence and Intellectual Property – Part II: Copyright and Artificial Intelligence. https://www.judiciary.senate.gov/imo/media/doc/2023-07-12_pm_-_testimony_-_sag.pdf Testimony before the U.S. Senate Committee on the Judiciary, Subcommittee on Intellectual Property.
- [45] J. E. Schirmer. 2020. *Artificial Intelligence and Legal Personality: Introducing “Teilrechtsfähigkeit”: A Partial Legal Status Made in Germany.* Springer, 123–142. https://doi.org/10.1007/978-3-030-32361-5_6
- [46] Eric Schwitzgebel. 2023. The Full Rights Dilemma for AI Systems of Debatable Moral Personhood. *ROBONOMICS: The Journal of the Automated Economy* 4 (2023), 32–32.
- [47] Eric Schwitzgebel and Mara Garza. 2020. Designing AI with rights, consciousness, self-respect, and freedom. (2020).
- [48] Y. Shi. 2022. On Negativism of Legal Personality of Artificial Intelligence. *Journal of Education, Humanities and Social Sciences* 1 (2022), 90–96.
- [49] B. F. Skinner. 2002. *Beyond Freedom and Dignity.* Hackett Publishing.
- [50] V. P. Talimonchik. 2021. The Prospects for the Recognition of the International Legal Personality of Artificial Intelligence. *Laws* 10, 4 (2021), 85.
- [51] Ted Wade. 2022. Artificial Persuasion: Ruled or Ruined by Self-Made, Unbounded, Amoral Advisors. *Substack* (26 February 2022). <https://tedwade.substack.com/p/artificial-persuasion> Accessed: December 1, 2023.
- [52] Zuzanna Warso and Paul Keller. 2023. Open Source, AI and the Paradox of Open. <https://openfuture.eu/blog/open-source-ai-and-the-paradox-of-open/>. Accessed: December 1, 2023.
- [53] Roman V Yampolskiy. 2012. Leakproofing the Singularity Artificial Intelligence Confinement Problem. <https://api.semanticscholar.org/CorpusID:13963004>
- [54] Eliezer Yudkowsky. 2000-2012. AI Box. SL4 Archive. <http://sl4.org/archive/0203/index.html3128> Accessed: December 1, 2023.
- [55] Eliezer S. Yudkowsky. 2023. The AI-Box Experiment. Singularity. <https://www.yudkowsky.net/singularity/aibox> Accessed: December 1, 2023.