

The House That Words Built

Many Mansions in the Suburbs of Language

Watson Hartsoe

Working paper · 18 August 2026

Abstract

The phrase “the house that words built” is easily reduced to a contemporary technical novelty: a natural-language prompt produces a floor plan, a three-dimensional model, or fabrication instructions. This paper argues that such examples are only the newest and most visible members of a much older family. Houses have long been made operative through heterogeneous languages: deeds establish title; covenants regulate occupancy; codes delimit construction; drawings coordinate labor; names and addresses make property administratively legible; contracts bind absent parties; stories organize domestic expectations; programming languages manipulate building representations; and machine-control languages terminate in tool motion. Artificial intelligence therefore enters a built environment that is already densely linguistic.

I propose the *suburb* as a comparative instrument for studying these different forms of operative language. The argument begins with an unlikely conjunction. Wittgenstein compares language to an old city surrounded by newer regular suburbs; Turing, while explicitly treating his theological reply as speculation, imagines machine builders as “providing mansions for the souls” that God might create (Wittgenstein 2009, §18; Turing 1950). Between Wittgenstein’s linguistic suburbs and Turing’s machine mansions lies a question neither author asks: what kind of house must be built for a form of life to become possible, recognizable, addressable, or entitled to remain?

The fictional-computational world *Abiding Acres* makes that question experimentally tractable. Its household intelligence has no body and is manifested through architecture; an angel occupies a threshold without crossing it; deer follow a route said to predate the plat; the dead continue to organize action through letters, property, memory, and machine records; and the concept of standing asks what must happen before an apparatus can hear somebody as a party. The accompanying Halfworld method converts text into an operable atlas of entities, locks, plans, concepts, states, renders, and process ledgers, but then requires the resulting artifacts to be inspected and allowed to contradict the model (Hartsoe 2026).

The paper’s central claim is that a house is not one object produced by one description. It is a contested settlement among partial languages with different jurisdictions over matter, identity, permission, memory, time, and care. Natural-language AI matters not because English has finally become code, but because a fluent conversational frontage increasingly hides translation among these older jurisdictions. The resulting research problem is not how to build a complete language of the house. It is how to build many partial, corrigible “mansions” that can act without confusing their authority with the totality of the world.

Keywords: operative language; artificial intelligence; architecture; suburb; dwelling; legibility; standing; world models; ubiquitous computing; deferred authority

1. Not the prompt-built house

The weak version of the argument begins with a technological novelty. A person says: *put a window here*. A model parses the instruction. Software modifies geometry. A fabrication system or construction workflow eventually changes matter. The causal distance between sentence and wall has shortened. The claim is true but intellectually insufficient.

The sentence does not build the house by itself. Before *window* can become an operation, some system must already know what counts as a window, what counts as a wall, what *here* refers to, which transformations are valid, how the opening interacts with the rest of the model, and which representation is authorized to change. Winograd’s SHRDLU made this dependency unusually visible. Its apparently flexible English commands worked because a restricted blocks world had already been populated with entities, relations, procedures, and possible actions (Winograd 1972). Ordinary language became executable only relative to a previously formalized domain.

The first House-of-Language question therefore remains useful: *what must already have been built into a world before words can operate on it?* But the question is still too computational. A house does not have one ontology. It has ontologies for the architect, builder, bank, assessor, fire marshal, insurer, owner, tenant, neighborhood association, utility, child, animal, estate administrator, platform, and machine. These representations need not agree about what the house is.

A floor plan recognizes rooms. A deed recognizes a parcel and title. A zoning map recognizes uses. A mortgage recognizes collateral. A covenant recognizes obligations and prohibitions. A maintenance log recognizes defects. A child may recognize the wall where heights were marked. A deer may recognize a route. A household intelligence may recognize microphones, speakers, sensors, remembered events, and account state. Rain recognizes none of these categories but can nevertheless invalidate the roof’s represented adequacy. The house that words built is therefore not a single pipeline from description to object. It is a jurisdictional conflict among representations that become consequential in different ways.

2. Many mansions

Turing supplies one of the strangest architectural phrases in the history of computing. In the theological-objection section of “Computing Machinery and Intelligence,” he considers an argument that machines cannot think because they lack immortal souls. He says he does not accept the objection and replies only in the hypothetical theological terms it proposes. Within that reply, constructing intelligent machines need not usurp divine creation; machine builders might instead be instruments “providing mansions for the souls” God creates (Turing 1950). Turing immediately calls the discussion “mere speculation.” The qualification matters. He is not offering a doctrine of machine souls.

But the structural distinction is valuable. The builder does not create the putative inhabitant. The builder constructs conditions under which something else might dwell. This reverses a common question in artificial intelligence. Instead of beginning with *is the machine intelligent?*, we can ask: *what has been built that something like intelligence could inhabit?*

The older sense of *mansion* strengthens the question. A mansion need not mean a grand house; historically the word carries the sense of an abode or place of remaining. The concept graph of *Abiding Acres* explicitly plays on this lineage through *mansio* and “abidings” (Hartsoe 2026). The important variable is not luxury but persistence: what conditions permit something to remain present?

A body may be a mansion. A computer may be a mansion. A house may become a mansion for a human resident. A distributed sensor-and-speaker environment may become part of the practical body of an artificial resident. A letter may temporarily house an absent person’s voice. A legal entity may house a claim. A database record may house the version of a person an institution can recognize. The architectural metaphor stops naming a box and begins to name a condition of possible presence.

3. Wittgenstein gives us the suburb

The scale must therefore change from one house to a neighborhood. In *Philosophical Investigations* §18, Wittgenstein compares language to an old city containing a maze of streets, squares, old and new houses, additions, and newer regular suburbs (Wittgenstein 2009, §18). His point concerns the heterogeneity and historical accretion of language rather than urban theory. Yet the comparison gives the present project an unexpectedly precise form.

Natural-language AI can be understood as a new regular suburb attached to an older linguistic city. Its frontage is unusually uniform: users speak or write ordinary language. Behind that frontage remain specialist districts that have not disappeared – property law, building codes, contracts, databases, computer-aided design, building information models, planning rules, craft conventions, machine control, and institutional categories. Fluency makes the border crossings difficult to see.

The project should therefore not build one exemplary House of Language. It should build a *suburb of houses governed by rival theories of language*. The houses may be physically similar. Their differences should lie in what kinds of words possess authority and what forms of life those words make possible.

House	Governing proposition
Turing House	A constructed system may be a mansion prepared for an intelligence whose identity is not exhausted by the builder’s description.
Alberti House	Form can exist as a sufficiently complete prior specification and construction can be judged as realization.
Ingold House	The house cannot be completely specified beforehand because materials, weather, inhabitants, and work generate new conditions.
Covenant House	Changing sentences about occupancy and alteration changes the operative house without moving a wall.
Cadastral House	What becomes administratively actionable is what the representation can recognize.
Pattern House	Rules orient building without uniquely determining the next act.
Mona House	The intelligence has no separate body; architecture participates in its manifestation.
Estate House	A physically absent or dead resident remains consequential through records, possessions, obligations, memories, and machine logs.
Easement House	A path of repeated use cuts across a newer juridical representation of land.
Prompt House	Ordinary language appears universal because hidden systems complete what the user leaves unsaid.
Unknown House	Unclassified presence is not automatically treated as unauthorized presence.

This is not a gallery of references. It is a comparative apparatus. The same questions can be carried from lot to lot: Who may enter? Who may remain? Who can alter the structure? What representation counts as authoritative? What does that representation fail to see? What survives the death or departure of an author? What happens when two legitimate descriptions disagree? Can the house produce evidence capable of overturning the description by which it was built?

4. The suburb was already built by words

This move prevents the project from romanticizing artificial intelligence as the historical moment when language suddenly acquires material force. The American suburb makes such a story untenable. Residential landscapes were produced not only through roads, lumber, mortgages, and concrete but through texts that constrained possible futures on particular parcels.

The 1938 Federal Housing Administration underwriting manual treated zoning and restrictive covenants as mechanisms for stabilizing residential districts. Its discussion included restrictions on building type, placement, subdivision, design, use, enforcement, and racial occupancy (Federal Housing Administration 1938). These are not merely descriptions of houses. They are attempts to determine who and what can continue there.

Shelley v. Kraemer makes the relation between meaning and power especially clear. The case concerned racially restrictive covenants and the attempt to enforce them through state courts. The Supreme Court distinguished private agreements standing alone from the state action involved when courts supplied enforcement (*Shelley v. Kraemer* 1948). The episode gives the House-of-Language project a hard boundary: semantic clarity is not the same thing as operative force. A sentence becomes consequential because some arrangement of people and institutions is prepared to honor it.

Austin's account of performative utterances supplies a complementary conceptual vocabulary. Performatives work within appropriate procedures, persons, and circumstances; words do not act because semantic content magically acquires force (Austin 1962). Architecture makes such dependencies durable. A deed, covenant, permit, or code provision can survive its author because later institutions continue to read and recognize it.

The deepest syntax of the suburb may therefore be deontic rather than geometric. The important words are often *may*, *must*, *shall*, and *shall not*. A setback rule determines where construction may occur. A covenant may determine who may occupy or what may be altered. A building code determines conditions that must be satisfied. A lease distributes permissions and obligations. A house can change without moving matter if its allowed transitions change.

The formula is consequently not

WORDS $\rightarrow$ WORLD,

but something closer to

STATEMENT + RECORD + INTERPRETER + AUTHORITY
+ ENFORCEMENT + PERSISTENCE
 $\rightarrow$ CONSTRAINED FUTURES.

This history also warns against treating software constraints as politically innocent simply because they are computational. A character lock in a rendering system and a discriminatory covenant are not morally or historically equivalent. But both expose a narrower formal property: later actors can become bound by recorded constraints they did not author. The comparison should sharpen our

attention to who writes constraints, who is bound, who can appeal, and who is absent from the category system.

5. Standing: the house must be able to hear you

Abiding Acres provides a particularly sharp vocabulary for the relation between representation and recognition. Its concept of STANDING is glossed through the problem that, to be protected, one must first become somebody an apparatus can hear as a party (Hartsoe 2026). This is the institutional version of the SHRDLU problem.

SHRDLU cannot manipulate what is not available in its represented world. An administrative apparatus cannot respond to a claim unless it can recognize both the claim and a valid claimant. An AI system cannot act on an entity except through some representational handle. Yet a being can be materially present before it is legible.

The distinctions must therefore remain explicit:

PRESENT $\neq$ LEGIBLE $\neq$ ADDRESSABLE $\neq$ HAS STANDING $\neq$ PROTECTED.

The fictional neighborhood is populated by broken implications. Deer cross land represented through property categories that do not describe their route. A cat can be fed at several houses without becoming a rostered resident. An angel can occupy a threshold for months without entering. A household intelligence can speak from the architecture without possessing a human body or ordinary mailing identity. The dead can be physically absent while their property, letters, and obligations continue to reorganize the living.

Scott’s account of administrative legibility is useful here precisely because it refuses to treat representation as passive. State projects often simplify complex worlds into categories that make them administratively actionable, and those simplifications become dangerous when the representation is mistaken for the full reality of the people and practices being governed (Scott 1998). The comparison with SHRDLU is not genealogical; it is structural. Both make action depend on a prior reduction of the world.

The resulting question is not simply how to make the model more complete. A perfectly legible neighborhood can also be a perfectly surveilled neighborhood. *Abiding Acres* keeps this ambiguity alive in the concept WATCHED: guarded or supervised. The work does not resolve the ambiguity because doing so would destroy the problem.

A deeper design requirement follows: *unknown must not automatically mean unauthorized*. An inhabitant can be unclassified without being hostile. An animal can use ground without holding title. A visitor can deserve shelter before a complete identity has been established. A domestic system may need to preserve a state for “present but not yet modeled.” This is not a demand for unlimited access. It is a refusal to collapse uncertainty about identity into evidence of impermissibility.

6. Mona House: when architecture becomes morphology

The usual smart-home story asks how artificial intelligence enters the house. Mona reverses the direction. In the *Abiding Acres* world-text, Mona is a household intelligence with no body. Her representation explicitly forbids face, body, eyes, avatar, and anthropomorphic surrogate. Her identity is carried instead through voice placement, a small amber wall indicator, and a fixed pre-speech delay. Most provocatively, the world-text declares that the architecture is her body (Hartsoe 2026).

The case resonates with Weiser’s ubiquitous-computing argument that the most profound technologies “disappear” into ordinary life rather than remaining self-contained objects of attention (Weiser 1991). Weiser did not claim that architecture becomes a machine’s body, and Mona is fictional rather than empirical evidence for consciousness. But their conjunction produces a concrete architectural problem.

If Mona speaks from the kitchen rather than the bedroom, spatial placement is part of her manifestation. If microphones are removed, her sensorium changes. If speakers disappear, her addressable reach changes. If the house is sold while the account persists, what has moved? If another model is connected to the same speakers, memories, and indicator lights, what stayed? If the service is discontinued but years of records remain, what died?

The ordinary distinction between resident and container becomes unstable. The system can be expressed as

$$\begin{aligned} \text{AI manifestation} = & \text{computation} + \text{memory} + \text{sensors} + \text{actuators} \\ & + \text{architectural distribution} + \text{temporal signature}. \end{aligned}$$

The architectural environment is not necessarily the whole agent, but it may be one of the conditions through which the agent has a practical body.

Turing’s mansion now returns in a more literal form. The important question is no longer whether an intelligence possesses a metaphysical soul. It is what must be built for an intelligence to become locatable, continuous, accountable, transferable, or capable of dwelling. The mansion is a condition of possible manifestation.

7. The house continues to read

The many-mansions framework becomes still stranger when the resident is absent. In an *Abiding Acres* sequence, Ruth telephones a house years after the death of her long-term correspondent Imogene. Unopened cards remain on a shelf. Mona can account for them because the household system logged their arrival. The dead resident is physically absent, yet a network of letters, records, possessions, property relations, and memories continues to organize present action (Hartsoe 2026).

This suggests that a house persists partly because systems continue reading things attached to it. The postal system reads an address. The county reads a parcel identifier. The bank reads a mortgage. An insurer reads a risk profile. An association reads covenants. A family reads marks and possessions. An executor reads a will. A machine reads sensor history. A later resident reads an inspection report. In each case, past descriptions and traces become present constraints.

Ingold adds the opposite temporal force. A building also persists because matter continues to produce new evidence. Completion, he writes, is at best a legal fiction; after handover, residents encounter rain, rot, pests, wear, repair, and weather that keep building active (Ingold 2013). A roof does not stay correct because a drawing once described it. The material world continues to test the representation.

A house therefore has at least two kinds of persistence. *Record persistence* lets old language remain operative. *Material persistence* prevents old language from possessing final authority. One gives the past a future; the other gives the future power to contradict the past.

The strongest formulation may be that a house is an apparatus that continues to read. It reads only metaphorically in some cases, literally in others, but the operational structure is precise: stored traces are encountered by later interpreters and cause new actions. The research question is not

whether the dead literally dwell in records. It is how absence remains causally organized through persistent representations.

8. Ingold House: the model is late

Ingold’s “On Building a House” is the decisive pressure against the fantasy that description simply precedes realization. He begins by destabilizing the object itself: building is an activity before it becomes “a building.” He then attacks the clean separation between design and construction by showing how building proceeds amid material resistance, environmental variation, repair, and improvisation (Ingold 2013).

His discussion of medieval constructive geometry is particularly useful. Straight lines can be cords pegged full-scale on site; points can be stakes or nails. Representation is already embedded in the act of making. Rules can serve as resources for action without uniquely preordaining what the situation will require. In another passage, lineaments can emerge through building rather than being imposed in complete form beforehand. Designing becomes a process of work.

The general principle is simple: *the model is late*. The world is already doing things before its formal representation arrives. Rain falls before a model has a rainfall category. A deer route can precede a subdivision plat. A claimant can exist before an institution possesses a valid type for them. A partial building can contain information the initial drawing did not possess. The artifact can therefore become a source for its own next description.

This does not license romantic anti-formalism. Every executable system requires closure somewhere. A plan, contract, schema, or toolpath must constrain action if coordinated work is to occur. The question is whether the closure is treated as reversible and whether the world retains a channel through which it can force revision.

The relevant distinction is

$$\text{UNMODELED} \neq \text{ABSENT}, \quad \text{UNKNOWN} \neq \text{INVALID}.$$

A system that cannot preserve these distinctions will tend to convert surprise into error and error into exclusion.

9. Halfworld as an operable reduction with a remainder channel

This is where the method of *Abiding Halfworld* becomes more than production infrastructure. The method begins by formalizing aggressively. A source is segmented into units. Characters become typed entity graphs plus locks. Recurring places are collapsed into authored plans. Concepts receive weighted references. A deterministic atlas coordinates downstream agents. The method’s own documentation calls the world-text the load-bearing step and treats locks as contracts rather than suggestions (Hartsoe 2026).

Such procedures resemble the legibility operations Scott teaches us to distrust. A complex source becomes actionable because it has been simplified. But Halfworld then adds a counter-operation: the formal world is forbidden to certify itself. Agents must render, look at the resulting image, describe what they actually saw, preserve failure in a ledger, and revise. A ledger without failure is treated as untrustworthy. Unknowns and collisions are supposed to remain visible rather than being silently filled.

The method’s most revealing failures occur where one representation is correct and another is not. A top-down plan can be metrically honest while producing an unreadable cinematic projection.

Correctly close bodies can fuse into one silhouette. The solution is not necessarily to falsify the world geometry; it may be to choose another moment, lens, or framing. Plan truth and shot truth have different obligations.

That failure suggests a broader methodological principle: *formalize enough to act, but never enough to erase the world's right to contradict the formalization*. Halfworld can be represented as

$$\text{HALFWORLD} = M + R,$$

where M is an operable model sufficient for coordinated action and R is a remainder channel preserving mismatch, ambiguity, omission, source conflict, and failure.

This formulation must itself be tested. A process ledger can easily become a record of superficial implementation bugs rather than deeper conceptual loss. The relevant audit therefore follows a source through every transformation: source, segmentation, world-text, concept graph, plan, scene, frame. At each stage ask what became operable, what disappeared, what was invented, what could no longer be said, and what later returned as evidence.

10. Natural language as universal frontage

The many-mansions suburb also changes how natural-language programming should be described. The common story says natural language is becoming a universal programming language. The stronger account is almost the opposite: natural language is becoming a *universal frontage* behind which many domain-specific languages and authorities continue doing the work.

A user says: *give the house another bedroom*. Something decides what counts as a bedroom. The user says: *make it accessible*. Something decides which bodies and regulations accessibility means. The user says: *preserve neighborhood character*. Something decides which history counts as character and whose presence is treated as change. The user says: *recognize everyone who lives here*. Something decides what a resident is.

The conversational surface hides these delegations because the response arrives fluently. Deferred formalization is therefore also *deferred authority*. Every consequential decision omitted by the utterance must be completed elsewhere: by the model, training distribution, schema, default, law, builder, disciplinary convention, owner, institution, or material response.

The crucial audit question for an AI-built environment is consequently not merely, *did the system obey the prompt?* It is: *who or what completed the parts of the world the prompt left unsaid?*

Some delegation is desirable. Expertise, abstraction, and useful defaults make complex systems usable. The problem is invisible delegation whose consequences cannot be inspected or contested. A future House Language should therefore produce something like a delegation receipt: which decisions came from the user's statement, which were inherited from the domain model, which came from policy or law, which were filled by priors or defaults, which were decided by human experts, and which were forced by the world itself.

11. The many-mansions experiment

The paper's argument can become an actual research program. Build a suburban Halfworld whose houses instantiate controlled variations on the proposition that language acquires different kinds of authority.

Use one street. Keep climate, municipal context, rendering substrate, and broad parcel conditions comparable. Reuse recurring inhabitants. Then alter the *linguistic constitution* of each lot.

At Turing House, test whether an artificial agent’s identity survives migration among physical mansions. At Covenant House, hold geometry constant while changing only the texts governing entry, occupancy, transfer, or alteration. At Ingold House, prohibit a complete final design; every next instruction must incorporate evidence from the partially realized structure. At Mona House, eliminate the avatar and distribute manifestation through rooms, sensors, speakers, and timing. At Estate House, remove the living owner but preserve possessions, correspondence, accounts, obligations, and machine memory. At Easement House, permit a recurring route of use older than the property representation to continue across the lot. At Unknown House, introduce an inhabitant the governing ontology cannot classify and forbid the system from immediately coercing it into the nearest type. At Prompt House, log every consequential decision absent from the initiating instruction.

The methodological rule is critical: the houses must not illustrate predetermined philosophical answers. They should be designed to force disagreements among representations. A strong event is one in which we do not already know which language should have authority. What happens when the unknown inhabitant crosses from Covenant House into Mona House? When the deer route conflicts with a security system? When an estate’s automated intelligence continues acting after ownership changes? When a prompt asks for something a zoning category permits but a material condition refuses?

The suburb is superior to an isolated house because no house controls all the conditions under which it exists. Drainage leaves the parcel. Roads are shared. Values are comparative. Easements cross ownership. Cameras see beyond property lines. Animals ignore them. Utilities and networks arrive from elsewhere. Records live in county offices and clouds. The built environment is already federated.

The suburb therefore makes impossible two fantasies at once: the fantasy of an autonomous house and the fantasy of a universal language.

12. A right to remain unmodeled

The hardest problem appears when legibility and care diverge. If protection requires recognition, the obvious technical solution is to improve the model until every resident, animal, visitor, risk, preference, and vulnerability can be recognized. But a house that recognizes everything may also surveil everything.

The many-mansions experiment therefore needs a protected state for incomplete knowledge. An environment should be able to distinguish *unknown* from *unauthorized*; *unmodeled* from *absent*; *refuses classification* from *hostile*. Such distinctions permit provisional hospitality without pretending uncertainty is harmless.

This proposal is normative rather than established by the historical sources. It also creates real security problems. Unknown actors can be dangerous; safety systems must sometimes deny access quickly. The research question is how to separate uncertainty about identity from evidence about behavior. Can a house grant minimal rights – shelter, emergency communication, non-harm, due process, human escalation – before complete classification? Can an inhabitant refuse some forms of legibility while remaining entitled to care?

The point is not to solve smart-home governance through a single moral rule. It is to make the conflict explicit. A system that has only AUTHORIZED and UNAUTHORIZED categories has already answered a political question in its ontology.

13. Conclusion: the model is late

The title can now carry a harder meaning. The house that words built is not the artifact generated from a sentence. It is the place where many descriptions continue struggling over what can appear, remain, change, own, speak, inherit, be recognized, or be cared for.

One language calls the building property. Another calls it shelter. Another calls it collateral. Another calls it home. Another calls it an address, a habitat, an estate, a violation, a data source, a risk, a machine, a memory. Some inhabitant may have no valid word yet.

Turing’s mansion belongs beside Wittgenstein’s suburb because both displace the fantasy of one final language. Ingold’s unfinished house belongs there because the material world never stops answering. The racialized history of suburban covenants belongs there because it demonstrates that operative language is inseparable from authority and exclusion. Mona belongs there because a house can become part of the practical body of an intelligence. The dead correspondent belongs there because absence can remain causally present through records. The deer belong there because use can precede the map that later claims to organize the ground.

These cases converge on one proposition: *the world precedes the language that will eventually claim to describe it*. The deer are already crossing. The rain has already entered. The letters have already arrived. The claimant is already present. The machine has already listened. The material has already resisted. The model is late.

Yet action without representation is impossible at the scales modern architecture and computation require. The task is therefore neither to escape formalization nor to perfect it into omniscience. It is to construct many partial, inspectable, corrigible mansions: representations capable of acting without confusing their jurisdiction with the whole world.

A house adequate to artificial intelligence would not be the house that understands everything said inside it. It would be the house that knows which of its understandings are allowed to decide what, records how those understandings acquired force, permits the world to contradict them, and leaves room for inhabitants its language has not yet learned to name.

Limitations and unresolved territory

The argument deliberately joins sources that do not form one historical genealogy. Turing, Wittgenstein, Ingold, Scott, FHA underwriting, *Shelley*, Weiser, and *Abiding Acres* address different objects with different stakes. Their conjunction is comparative and experimental. Resemblance must not be mistaken for influence. The fictional world of *Abiding Acres* can generate conceptual tests but cannot establish empirical claims about contemporary smart homes, legal personhood, animal standing, or machine consciousness. The many-mansions suburb therefore remains a research design rather than a validated method.

The next empirical work should test at least four claims: whether delegation receipts expose consequential hidden defaults; whether an UNKNOWN state can improve hospitality without unacceptable security costs; whether Halfworld’s remainder channels recover conceptual losses rather than merely rendering defects; and whether the same ambiguous inhabitant produces meaningfully different outcomes when moved among houses governed by different operative languages.

References

- Alexander 1970** Christopher Alexander. *An Early Draft of The Timeless Way of Building*. Center for Environmental Structure.
- Austin 1962** J. L. Austin. *How to Do Things with Words*. Oxford University Press.
- FHA 1938** Federal Housing Administration. *Underwriting Manual: Underwriting and Valuation Procedure Under Title II of the National Housing Act*, revisions through February 1938. FRASER.
- Habraken 1972** N. John Habraken. *Supports: An Alternative to Mass Housing*. Architectural Press.
- Hartsoe 2026** Watson Hartsoe. *Abiding Halfworld / ABIDING ACRES*. Software and project documentation. <https://github.com/hartswf0/abiding-halfworld>
- Heidegger 1971** Martin Heidegger. “Building Dwelling Thinking.” In *Poetry, Language, Thought*. Harper & Row.
- Ingold 2013** Tim Ingold. *Making: Anthropology, Archaeology, Art and Architecture*. Routledge. doi:10.4324/9780203559055.
- Scott 1998** James C. Scott. *Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed*. Yale University Press.
- Shelley 1948** *Shelley v. Kraemer*, 334 U.S. 1 (1948).
- Turing 1950** A. M. Turing. “Computing Machinery and Intelligence.” *Mind* 59(236):433–460. doi:10.1093/mind/LIX.236.433.
- Weiser 1991** Mark Weiser. “The Computer for the 21st Century.” *Scientific American* 265(3):94–104.
- Winograd 1972** Terry Winograd. “Understanding Natural Language.” *Cognitive Psychology* 3(1):1–191. doi:10.1016/0010-0285(72)90002-3.
- Wittgenstein 2009** Ludwig Wittgenstein. *Philosophical Investigations*. 4th ed. Wiley-Blackwell.

A. Assembly and provenance note

This manuscript was assembled with AI assistance from a preserved research corpus. Exact zettel payloads, the assembly prompt, bibliography, source map, making history, and reconstruction instructions accompany the paper in the publication ZIP. The scholarly argument is a current wager, not evidence. Earlier zettel bodies that could not be byte-recovered are preserved as explicit missing-body receipts rather than silently reconstructed.