

The Noise of Sculptors

Bugs, Finished Sketches, and Creative Judgment in the First Midjourney Discord

Watson Hartsoe

Working Paper - 17 August 2026

Abstract

The first thing Midjourney taught me was not how to prompt. It was how quickly an image could disappear. In the fall of 2022, entering an image-generation channel meant entering a flood: prompts, four-image grids, variations, jokes, failed experiments, professional concept art, grotesque accidents, copied phrases, and strangers talking past one another while the screen kept moving. Our field notes called the place “chaotic,” “messy,” and “fast.” A participant remembered that work could be “washed away immediately”; in busy rooms, more than one hundred new images could appear in ten minutes (Sun, Hartsoe, and Ottolin 2022, 42). We later gave this world a neat title: *Sculptors of Noise*. The title now feels backward. Before anyone could sculpt the noise, there was the noise of sculptors.

This paper returns to that noise through a small sequence preserved unevenly across screenshots, field notes, the 2022 study, and a later private exchange. During a weekly Midjourney office hour, I asked a question about originality and ceding creative control. A user named Joshua responded that the most inventive styles he had found came from “bugs and quirks.” He later told me about a phrase that generated an architectural style he could not explain, solved a design problem in a VR game, and became the center of more than ten thousand experiments. The most consequential part of his account was not the magical prompt. It was his explanation of why the design persuaded him: “the only way it sold me is because it was finished.”

That sentence changes the creative problem. Conventional workflows often require creators to decide which ideas deserve expensive refinement while those ideas are still crude. Generative systems can reverse that sequence: render first, judge later. They make it cheap to bring unlikely ideas to premature finish. The result is not frictionless creativity. It is a new selection problem inside abundance. The early Midjourney Discord made that problem public: almost everything could look finished enough to matter, while almost nothing remained still long enough to understand. The central craft therefore shifts from making without error toward learning which errors, deviations, and weak-looking ideas deserve not to be corrected yet.

Keywords: Midjourney; generative AI; creativity; creative judgment; prompting; design; iteration; Discord; error; selection

The first channel

The first thing I remember is losing the image.

I had typed a prompt into one of Midjourney’s Discord generation channels. The bot answered, but the answer did not stay where I had left it. Other people were prompting at the same time.

Their images arrived above and below mine. Four-image grids stacked into one another. Someone asked for a dragon. Someone else asked for a building. Portraits appeared, then disappeared upward. Variations arrived before I had finished looking at the first result.

The channel did not feel like a studio. It felt like standing beside a conveyor belt that had learned to dream.

You could scroll back. That was not the point. The point was that the interface taught a bodily lesson before it taught a technical one: **there was too much to see.**

Our original paper preserved a cleaner version of this sensation. We wrote that when we entered the community we found it “chaotic,” “messy,” and “fast.” Hematita remembered that in the channels for sharing work, an image could get “washed away immediately.” In the generation rooms, more than one hundred new images could appear within ten minutes (Sun, Hartsoe, and Ottolin 2022, 42). People later made private threads so their own work would not vanish inside the “huge wave” of everyone else’s images. Advanced users complained about “walls of images” and moved into more exclusive rooms. The interface did not merely display abundance. It continuously erased the local importance of whatever had just happened.

That was the research site.

The problem with writing about it four years later is that the mess now wants to become a theory. We know the words that came afterward: prompt engineering, human-AI collaboration, latent space, control, co-creation, creative support, generative workflow. Each phrase makes the scene easier to discuss and harder to remember.

The scene did not arrive as a concept.

It arrived as too many pictures.

We eventually called the paper *Sculptors of Noise*. Clarinet had given us the phrase. He said Midjourney users sculpted from noise the way another artist might sculpt from marble; prompts were chisels with force, angle, strength, specificity, and other ineffable properties (Sun, Hartsoe, and Ottolin 2022, 1). It was a good line. It defended human craft at a moment when people were already asking whether the machine had made the art by itself.

But the line begins after the difficult part.

“Sculptors of noise” imagines a person holding a tool against a material. The sculptor remains singular. The material waits. The metaphor gives the human a hand, an instrument, and an object of work.

Discord gave us none of that calm.

There were thousands of hands. The tools changed. The material answered back. Other people’s experiments entered your field of vision whether you wanted them or not. A useful phrase could cross the server before anyone knew why it worked. A bad result could be a mistake, a clue, a joke, a failure, or the beginning of somebody else’s project. A model update could make yesterday’s technique stop behaving like yesterday’s technique. Moderators could close one line of inquiry. Super-user rooms could hide another one. Professional users could hoard prompts.

Newcomers could arrive expecting customer service while early users remembered something closer to people gathering around a newly discovered fire.

The phrase I want now is the reversal:

the noise of sculptors.

The noise was not merely volume. It was unfinished knowledge: prompts, images, guesses, rituals, workarounds, arguments, private rooms, spreadsheets, conspiracies, repeated tests, screenshots, and stories about what the machine “liked.” Knowledge did not sit above this mess waiting to be extracted. Knowledge grew inside it.

The question on stage

One of the field’s most consequential moments survives in pieces rather than a clean transcript.

A screenshot from a Midjourney Office Hours shows the Discord stage. Three people are speaking. Seven hundred and sixteen are listening. DavidH is on stage. So am I, under the username **cool radio**.




Figure 1. Midjourney Office Hours in Discord. The field account “cool radio” is on stage beside DavidH; the interface shows 716 listeners. User-provided field archive, 2022.

I remember asking a question after David had spoken about Chinese philosophy. My notes from the period circle the same problem again and again: *To be creative, should we abandon the search for originality? Is creativity an act of submission? Is creativity an act of control? Is imagination a form of humility or ego? Would a purposefully boring approach yield the most unique aesthetic?*

Those notes were shaped by an essay by philosopher Julianne Chung on a reading of the *Zhuangzi*. Chung contrasts the familiar image of creativity as originality with a model of creativity as integration: a craftsperson responds sensitively to the particular situation rather than imposing a complete plan. Her Wheelwright Pian example emphasizes a kind of practical knowing that cannot be reduced to rote instructions. She also makes a claim that mattered to me in 2022: aiming directly at originality can narrow search, because we explore only the regions we already expect to look original (Chung 2020).

I do not have evidence for the exact sentence I spoke on stage. The screenshot does not contain it. Memory has had four years to make the question cleaner. I do have two independent traces of its substance. First, my notes preserve the originality/integration problem. Second, our 2022 paper says that Joshua contacted the research team after we asked during office hours how best to “cede creative control to the AI to encourage original image generation” (Sun, Hartsoe, and Ottolin 2022, 14).

That is enough to preserve the question without inventing a transcript.

The important thing is what happened next.

Joshua was listening.

A second screenshot preserves his response in the Discord discussion:




Figure 2. Joshua’s written response after the office-hours question: his most inventive styles had come from “bugs and quirks.” User-provided field archive, 2022.

“Great question. Especially considering that the most inventive styles I’ve found have come from essentially bugs and quirks”

I thought I was asking a philosophical question about originality.

Joshua answered with error.

The answer that left the stage

Later that day, I sent Joshua a direct message.

“What types of bugs have you found that produce creativity?”

He answered that the most common involved misinterpreting artist and cultural references. Because the interpretation was wrong, it could also become more unique. He had theories about why. He was not sure.

“yeah,” he wrote. “tricky!”

Then he asked what I was studying.

The answer I gave him is more useful to me now than much of the polished language I later placed around it. I said I was interested in creative practices that could not be fully captured by an algorithmic list of instructions. I wanted to know how people responded to the precise particularities of AI output and interface affordances through previous knowledge, making what they wanted not by imposing a complete plan but by adapting and “going with the flow” to some extent.

I was borrowing the shape of Chung’s Wheelwright Pian. The important knowledge was supposed to live in the adjustment: not too hard, not too soft; not pure plan, not pure accident.

Joshua gave that abstraction a body.

He had been using Midjourney since June. He used it mostly for concept design on a VR game. He estimated that he had made about 24,000 images.

I asked whether those thousands of images represented thousands of concepts or whether he worked one concept again and again.

His answer broke the research open.

He had found a phrase that seemed to invent a new architectural style. He did not know why. He guessed that perhaps a shape became washed out early in generation, Midjourney reinterpreted the residue as a surface, and then “a bunch of conclusions follow.” Then he stopped himself: “Dunno tho.”

That uncertainty did not weaken the discovery.

The phrase solved a real design problem.

Joshua started rebuilding the game’s art direction around it. More than ten thousand of his images came from that phrase and combinations around it. He called it a “magical prompt.” He had found it in his second week using Midjourney.

“After that everything changed hehe.”

The obvious story is that an AI system helped invent a novel style.

The more important story begins where novelty stops.

The idea that would have died as a sketch

I asked Joshua what had been so convincing about the new direction.

The studio had earlier art that vaguely resembled it. They had already moved away from that direction because it did not seem good enough. The idea had been seen, in other words. It had been judged.

Then Midjourney returned something close to that discarded territory in a more resolved form.

Joshua wrote:

“the only way it sold me is because it was finished.”

That sentence is the paper.

Not “AI expands the possibility space.”

Not “the user gives up control.”

Not “prompting is system identification.”

Those claims can all be useful. None reaches the nerve.

The machine changed when judgment happened.

A conventional design process often begins with cheap representations. A thumbnail costs less than a polished render. A blockout costs less than a finished environment. A storyboard costs less than a shot. The cheap representation lets the creator kill weak ideas before paying to finish them.

That economy hides an assumption: the rough representation tells us enough about the finished idea to decide whether the idea deserves refinement.

Joshua found an exception large enough to reorganize a professional project.

The rough version had lied.

Or, more precisely, it had told the truth at the wrong resolution.

The direction looked weak while crude. It looked persuasive when pushed farther toward finish. Under a conventional workflow, abandoning it had been rational. Why spend scarce labor finishing an idea that already appears unlikely to work?

Joshua made the contradiction explicit. There exists, he argued, a “giant set of designs” that do not seem viable from the side of “thumbnail sketches” and become discoverable only through more finished exploration. Most people do not have time to carry all those questionable ideas to a finished state. Even if they did, the act seems absurd: **why finish what you already think will not work?**

Generative AI changes that calculation.

It makes the absurd experiment cheap.

That is more consequential than “faster ideation.” Faster ideation leaves the old order intact:

imagine → sketch → judge → refine.

Joshua describes another order:

imagine → render → then decide whether the idea was worth imagining.

The distinction sounds small until we notice what traditional pipelines discard.

They do not discard only bad ideas.

They discard ideas whose value appears too late.

A thumbnail cannot prove atmosphere. A wireframe cannot prove delight. A blockout cannot prove the emotional force of a building. An ugly paragraph cannot prove the book hiding inside it. Professionals learn to see through impoverished representations. That skill is part of taste and judgment. Nobody sees through perfectly.

Some ideas die because they look stupid too early.

Some strange ideas need more finish before anybody can tell they are alive.

Midjourney changed the price of finding out.

When not to correct the machine

Joshua's first answer now comes back with greater force: "bugs and quirks."

The normal account of expertise assumes that a skilled practitioner detects error faster and corrects it sooner. Joshua describes a craft that sometimes does the opposite.

The system misreads a reference.

A shape washes out.

A surface appears where no surface was requested.

A style drifts away from the intended example.

The obvious move is correction:

No. Not that. Try again.

But the expert sometimes makes another move:

wait.

What if the mistake knows somewhere I do not?

This is the point at which the question about control becomes useful again. Submission is the wrong word. Joshua did not hand the game to Midjourney. He generated ten thousand images, judged them, discarded branches, combined effects, and redirected the project. Control is also the wrong first word. The creative event began before control returned.

It began when he allowed a deviation to remain alive.

Creativity can begin by postponing correction.

That does not mean mistakes are inherently creative. Most errors are dead ends. Noise is mostly noise. A craftsperson who follows every deviation has not become humble; they have stopped choosing.

The skill lies in discrimination.

Which error deserves four more generations?

Which one deserves forty?

Which strange result deserves ten thousand?

This is where the Wheelwright Pian analogy becomes more than decoration. In Chung's account, the wheelwright's knowledge lies in a responsive fit between person, material, tool, and situation rather than in the execution of a complete abstract plan. The Midjourney practitioner faces a very different material, but the practical tension rhymes: impose too hard and the system

becomes a bad executor of an already known idea; surrender completely and the result loses purpose. The craft lives in the adjustment.

The important difference is that Midjourney can generate deviations at industrial speed.

The practitioner must therefore develop not only sensitivity to material but **stopping rules for surprise**.

Finished too soon

Now the first Discord channel looks different.

It was full of things that had reached visual finish before they had earned social importance.

A novice could type a few words and receive something with lighting, texture, composition, and atmospheric authority. A professional could generate dozens of plausible directions before lunch. A joke could arrive with the finish of an illustration. A dead end could look like a poster. A mistake could look expensive.

The old visual cue that said **somebody spent time on this** had begun to separate from the image.

That separation created the flood.

If almost every experiment can arrive with polish, polish stops filtering attention. The channel fills with premature finish. The user gains the power Joshua valued - the power to inspect unlikely ideas at a later stage of representation - and immediately inherits the opposite problem: **too many things now look worth inspecting**.

The abundance and the insight are the same event seen from different sides.

Joshua experienced it as access to designs that rough sketches would have killed.

A newcomer experienced it as their own image disappearing under one hundred others.

The platform lowered the cost of visual completion without lowering the cost of human attention.

That mismatch is the noise.

Generative AI therefore moves a bottleneck. Before, finish was expensive enough to ration exploration. After, exploration can become cheap enough that attention, memory, and judgment become the scarce materials.

The community responded by building dams.

People made private threads so their work would not disappear. Hematita built a smaller server because the main one moved too quickly. High-volume users entered gated rooms where “arcane knowledge” supposedly circulated. Some of those rooms filled up too and were described as “gentrified.” Joseph kept professionally valuable prompts secret. Cody Boy called prompts “trade secrets.” MariusJuston built an artist-style spreadsheet. Ignite built randomized prompt protocols. The public feed created abundance; expertise often required constructing a smaller world inside it.

The magical prompt was not a spell

The phrase “magical prompt” tempts the wrong reading.

It sounds like superstition: secret words that possess occult power over the model. Our fieldwork contained plenty of that language. Users talked about magic words, prompt leakage, ghosts in the machine, mountain peaks, valleys, “fuzzy” models, and “thicker” models. Breakthroughs came from glitches and happy mistakes as well as quasi-experimental trial and error (Sun, Hartsoe, and Ottolin 2022, 47-50).

Joshua’s phrase matters because it reveals what the magic did in practice.

The phrase did not produce one astonishing image and end the search.

It produced a neighborhood worth returning to.

Joshua tried styles. He hit dead ends. When he saw progress, he iterated and honed in. More than ten thousand images accumulated around a phrase he could not causally explain.

The magic was not certainty.

The magic was **yield**.

A useful prompt produced enough consequential variation to sustain further attention.

This also changes how we should read the community’s folk theories. Mountains, valleys, magic words, and ghosts were weak substitutes for a mechanistic account of a proprietary model. They were also practical devices for deciding what experiment to run next. A metaphor can be wrong about the machine and useful about the next move.

The line between craft and superstition therefore does not run cleanly between true and false explanation. It also runs through experimental consequence.

Can the belief produce a discriminating test?

Can another user reproduce the effect?

Does the effect survive a new seed?

Does it survive a new model version?

Does it open a productive neighborhood or merely decorate one lucky output?

The community learned inside those questions because the system rarely supplied the answers.

The public rough draft

Midjourney made another choice that now feels difficult to recover as strange: much of the act of making happened in public.

Oscar gave us the cleanest description. Imagine using Photoshop with a setting that automatically published every autosave to your website and pushed each save onto social media. Nobody would design a normal creative tool that way. In Midjourney, something like it was the default social condition (Sun, Hartsoe, and Ottolin 2022, 45-46).

That design created a new public object: **the rough draft that looks finished.**

Traditional creative communities often see selected final work and retrospective explanation. The failed thumbnail stays on the desk. The abandoned layer stays in the file. The embarrassing intermediate version never enters the gallery.

Midjourney inverted that privacy. Prompt, output, variation, joke, failure, revision, and breakthrough could all flow through the same public channel.

That made the community an accidental school.

Oscar said that seeing other people's work and remixing pieces of it compelled him to improve. But the same publicness made professional secrecy rational. Joseph hid the prompt around his game. Cody Boy called prompts trade secrets. Clarinet described his images being scraped and sold. Public process produced apprenticeship and appropriation at the same time.

The early Discord therefore did not merely host a tool.

It exposed the unstable politics of a new creative resource: **the path by which a finished-looking artifact came to exist.**

When production becomes abundant, the path can become more valuable than the artifact.

The archive we should have kept

We saved interviews.

We saved screenshots.

We saved quotes.

I wish we had saved more failures.

Joshua's final prompt cannot show us the craft. The successful images cannot show us the craft. Somewhere among those ten thousand generations were the dead ends, the almos, the mutations, the points where a useless direction became useful, and the moment when an output changed what the creator wanted.

That path was the practice.

If I were designing the study again, I would not begin by asking for the best prompts. I would ask for a prompt genealogy.

Show me the first attempt.

Show me what disappointed you.

Show me what you changed.

Show me the mistake you nearly deleted.

Show me the branch you followed too long.

Show me the moment your goal changed.

Show me what stopped working after the model changed.

Then show me the final artifact.

The field's most durable evidence would not be the magic phrase. It would be the sequence of judgments that allowed one deviation to survive while thousands of others died.

This also answers the question I was trying to ask on stage.

Should creativity abandon originality?

The question was too clean.

Originality is not something a creator either pursues or renounces. It can be an outcome of how the creator distributes attention across familiar intentions, strange errors, other people's work, and unexpected consequences. Direct pursuit can narrow the search. Total surrender can dissolve the project. What matters is whether the practitioner can keep a surprising possibility alive without confusing surprise with value.

The machine contributes variation.

The human contributes a changing threshold for commitment.

Neither owns creativity alone.

What the noise teaches

The first Midjourney channel looked like excess because it was excess.

Too many images.

Too much finish.

Too many guesses.

Too many people discovering the same trick, inventing different explanations, hiding useful phrases, copying each other, testing boundaries, and watching the model change under their hands.

We tried to make that legible by writing about control.

Control was not wrong.

It was late.

Before control comes a more fragile act: **noticing something that your current plan would have corrected away.**

Joshua's architectural style began there. A bug or quirk produced something he did not ask for. A rough cousin of the idea had already lost in the conventional process. Midjourney let the idea return farther downstream, dressed in enough finish to argue for itself. Joshua then spent ten thousand images deciding whether the argument held.

The machine did not simply answer him.

It changed what was worth asking for.

That is why I now prefer *The Noise of Sculptors* to *Sculptors of Noise*.

The old title begins with mastery.

The new title begins with the teeming field in which mastery has not yet learned its own name.

That is closer to what the Discord felt like in 2022. It is also closer to the creative problem generative systems leave us with now.

When finish becomes cheap, judgment has to become more careful.

When error becomes plentiful, correction cannot remain automatic.

When possibility becomes abundant, the creative act is no longer simply to make more.

It is to learn **what not to throw away**.

Limitations and unresolved territory

This paper deliberately treats a small historical archive as a source of unresolved questions rather than as a representative survey of generative-AI practice. The original study involved more than one hundred hours of participant observation and twelve interviews, but the Joshua sequence remains one professional case. His account cannot establish how often finished exploration rescues ideas that thumbnails would have killed.

The office-hours origin is also incomplete. The surviving stage screenshot establishes the setting, audience scale, and my presence on stage. Joshua's screenshot establishes his "bugs and quirks" response. The 2022 paper records the substance of the question as ceding creative control to encourage original generation. My contemporaneous notes preserve the originality/integration problem. None preserves the verbatim spoken question. I therefore refuse to quote a sentence that the evidence does not contain.

Finally, Joshua's causal story about shapes washing out during generation was explicitly speculative. It matters as a practitioner hypothesis, not as a verified description of Midjourney's internal architecture. The strongest claim does not depend on that mechanism. It depends on the reported change in creative judgment: an earlier direction looked unpromising until a more finished representation made it persuasive.

Those boundaries leave the most important work open. We need experiments that compare judgments across representation fidelity. We need process archives that preserve failures rather than only successful prompts. We need to know whether experts differ from novices in which deviations they keep alive. And we need to know what happens when generative abundance produces more apparently finished candidates than anyone can inspect.

The first channel already contained all four problems.

We simply did not know their names yet.

References

Chung, Julianne. 2020. "To be creative, Chinese philosophy teaches us to abandon 'originality'." *Psyche*, September 1. <https://psyche.co/ideas/to-be-creative-chinese-philosophy-teaches-us-to-abandon-originality>

Sun, Zhoujun, Watson Hartsoe, and Tommy Ottolin. 2022. *Sculptors of Noise: Control in Midjourney AI Art Community*. CS 6470: Design of Online Communities, Georgia Institute of Technology. User-provided working paper preserved with this research field.

Appendix A. Evidence path for the office-hours sequence

The paper keeps four evidence objects distinct:

1. **Office-hours screenshot.** Shows the Discord stage, “cool radio,” DavidH, three speakers, and 716 listeners. It does not show the spoken question. 2. **Joshua response screenshot.** Shows Joshua’s written response that his most inventive styles came from “bugs and quirks.” 3. **2022 research paper.** States that Joshua contacted the researchers after an office-hours question about ceding creative control to encourage original image generation. 4. **User-provided 2026 field recollection and surviving notes.** Connect the question to Chung’s essay, originality versus integration, submission versus control, and the possibility that direct pursuit of originality narrows creative search.

The reconstruction treats agreement among these traces as support for the question’s conceptual neighborhood, not permission to manufacture a verbatim line.