Executable Repertoire Before the Canon: Early Generative-AI Practice as Social Learning

Research field working paper

17 August 2026

Abstract

This working paper argues that an important unit of early generative-AI expertise was neither the isolated prompt nor accurate knowledge of model internals, but a socially circulated repertoire of executable experiments. A small body of early Midjourney material associated with animator and educator Scott Roberts is useful because it spans informal collaborative play, explicit improvisational vocabulary, and an experimental university course in which students were described as research collaborators. Read alongside situated-learning and distributed-cognition theory, the field suggests three propositions: prompt sharing can make repertoire executable before it is understood; newly born technical communities can produce relative "old-timers" before they possess a canon; and a generative model can enter the cognitive organization of collective practice without thereby becoming a social member of the community. These are current wagers, not settled historical claims. Three key uncertainties remain: the transcript's "Scott" has not been independently tied to Roberts; Roberts's exquisite-corpse posts have not been shown to be Yelmo; and temporal proximity between 2022 experiments and 2023 pedagogy does not prove transmission.

1 A practice before a canon

The supplied 2022 prompting notes describe a community learning by watching other people's prompts, comparing outputs, withholding or sharing hard-won formulations, and developing explanations for why certain words seemed to work. The notes repeatedly position expertise as something acquired through comparative practice: users see work better than their own, inspect the prompts behind it when those prompts are shared, experiment, and develop local theories about weighting and training data. The archive therefore already resists a simple equation of expertise with knowledge of model internals. Card `PROMPT-FORAGE-001` follows that edge into Jonas Oppenlaender's taxonomy of prompt modifiers, where prompt engineering is treated as iterative probing and community practice becomes analytically visible.

This distinction matters because operational success and causal understanding can separate. A prompt term can repeatedly produce a desired effect while the practitioner's explanation of that effect is wrong. Card `PROMPT-FORAGE-002`, for example, opposes a folk retrieval story – a word "references a set of data" at prompt time – to published latent-diffusion machinery in which text conditions generation through learned representations and attention. The field does not use that technical paper to reverse-engineer Midjourney. It uses it to establish a methodological rule: observed prompt efficacy does not validate the explanation practitioners attach to it.

Etienne Wenger's account of communities of practice helps name the social dimension, but the prompt adds a peculiar material twist. A shared craft repertoire is ordinarily demonstrated, narrated, copied by hand, or embodied in tools and routines. A prompt can be pasted into the generative system and made to produce another observable event immediately. The novice can therefore execute a fragment of somebody else's practice before understanding it. Card `COP-FORAGE-001` calls this *executable repertoire*. The phrase is a synthesis of this research field, not Wenger's terminology.

Executable repertoire changes the order of apprenticeship:

shared artifact → run → observe consequence → modify → partial understanding.

Visible competence can consequently precede causal understanding. Copyability can accelerate entry while also carrying superstition, model-version artifacts, and selectively remembered success. A prompt community can accumulate genuine operational knowledge and inaccurate explanations at the same time, and the same shared prompt can transmit both.

2 Improvisation as a social protocol

Scott Roberts's public August 2022 material supplies a more specific form of circulation. In one post he described "an on-going exquisite corpse game" conducted through Midjourney; a collaborator independently described an "exquisite corpse collaboration." Card `SCOTT-COP-001` treats that language as evidence of a social protocol rather than merely an aesthetic label.

This changes how the "riffing" described in the supplied notes can be read. The notes describe one participant proposing a Muppet, another extending the idea into space, and another moving toward a Muppet space pirate. The activity is compared to a musical experience because many people can respond in real time without literally painting over one another's work. In an exquisite-corpse-like process, no participant needs to author the whole object. A turn inherits a residue, transforms it, and passes a changed situation onward.

The generative model sits inside that chain as a transformation device whose unexpected result can become the next person's provocation. The prompt may therefore be the wrong fundamental unit. The stronger unit is a *turn*: an inheritance, a human intervention, a machine transformation, and a socially visible consequence.

That move also alters the genealogy of "AI creativity." Roberts contributed a chapter titled *Surrealist Storytelling: Artistic Experimentation and Spontaneity through Improvisation*. Card `SCOTT-COP-002` does not treat title similarity as proof of influence. Instead it opens a historical test. Generative AI's celebrated surprise may, for some practitioners, be less a novel creative method than a new instrument for an older method that already values interruption, incompletion, contingency, and response.

The responsible claim is narrow: the same practitioner appears in public records joining Midjourney to exquisite corpse and in publisher metadata joining his work to Surrealist improvisation. Whether that is one continuous practice genealogy remains unresolved until earlier Roberts exercises, teaching materials, or the full chapter can be recovered.

3 Old-timers measured in weeks

Lave and Wenger reconceive learning as movement through participation in a community of practitioners. Their newcomer/old-timer relation becomes strange when the practice itself has only recently appeared. In a newly released generative platform, differential exposure can manufacture relative expertise at extraordinary speed. Someone generating daily for six weeks can possess a repertoire inaccessible to a person arriving today, even though neither participant has long historical experience of the medium.

Card `COP-FORAGE-002` therefore asks whether seniority can be generated faster than canon. The community may possess people recognized as advanced before it possesses stable explanations, textbooks, standards, or even a durable version of the tool. Major model changes can also destroy some of that advantage overnight.

Scott Roberts complicates any purely platform-native hierarchy. The user-supplied biography and DePaul materials place him within long-standing animation, visual-development, design, teaching, and improvisational practice. Platform tenure and transferable craft therefore cross-cut one another. A student may discover yesterday's new feature before the professor does while the professor retains much deeper expertise in animation continuity, critique, staging, production, and visual development. "Expert" becomes local to the problem being solved rather than a permanent rank.

Roberts's later pedagogical language makes this structure especially useful. Public materials surrounding DePaul's experimental AI Art Tools course call students "student research collaborators," and Roberts publicly celebrated an "AI Art Tools research team" that had "figured out" a tool. Card `SCOTT-COP-003` treats those phrases as a testable claim about knowledge flow, not proof of egalitarian pedagogy. The archive needed to establish the claim would show techniques moving student-to-teacher and peer-to-peer, not only teacher-to-student.

The deeper possibility is that the course was not downstream from a stabilized craft. It was one of the sites where the craft was being stabilized.

4 From game to curriculum: the missing middle

The chronology is compressed. Roberts publicly described Midjourney exquisite-corpse activity in August 2022. DePaul records a May 5, 2023 presentation of "The Cyborg Muse," and Roberts's faculty biography states that in 2023 he created a course devoted to generative-AI art tools for animation, film, and theater. Card `SCOTT-COP-005` preserves this temporal proximity while refusing to call it transmission.

The strongest research object is therefore the missing middle. Did practices such as riffing, exquisite corpse, remix, character consistency, shared prompt experimentation, and tolerance of machine surprise migrate into the course? Or did the classroom develop a substantially different repertoire around a rapidly expanding tool ecology?

That question can be answered historically. A course proposal, dated syllabus versions, chronological assignment sheets, student projects, Discord or Slack messages, and teaching presentations would allow each technique to be traced from first social occurrence to first pedagogical occurrence. Without that evidence, "practice became curriculum" is a hypothesis with a dateable gap.

This missing middle is methodologically preferable to a smoother story. It gives the research something capable of dying.

5 The model inside cognition, outside membership

Roberts's "Cyborg Muse" material introduces another recursion. The proposal describing human-machine creativity disclosed that it was itself created "in collaboration with ChatGPT." Card `SCOTT-COP-004` follows that fact upward: the technology under investigation can participate in producing the language by which the community explains and institutionalizes its investigation.

One response is to ask whether the AI is a collaborator or merely a tool. Edwin Hutchins's distributed-cognition framework suggests a more discriminating question. Cognitive work can be organized across people and representational artifacts without requiring every causally important component to have human-like social standing.

Card `COP-FORAGE-003` therefore separates *cognitive participation* from *social membership*. In collaborative Midjourney play, a generated image can redirect attention, materialize a possibility no participant explicitly proposed, and coordinate the group's next move. The generator can be inside the causal and representational organization of the activity. None of that establishes intention, authorship parity, accountability, understanding, or membership in the human community of practice.

This distinction avoids two opposite reductions. The generator need not be inert merely because it is not a person. But causal contribution does not automatically make it a social peer.

The test is move-level reconstruction. Before each generation, record what human participants already imagine; after the output, record possibilities that appear only because of what they saw. Compare this with a condition in which humans receive shuffled or randomly selected human-made images. The empirical question is not whether the machine "is creative" but where new representational state enters the distributed activity.

6 Executable apprenticeship

Taken together, the field supports a provisional model of *executable apprenticeship*. Early participants encounter an opaque and changing generator. They learn by comparing artifacts, sharing prompts, remixing one another's moves, and circulating explanations. Some of the repertoire is directly runnable. Machine outputs become both demonstrations and provocations. Relative expertise differentiates quickly because experiment rates and prior craft are uneven. Formal pedagogy can begin before stable technical consensus exists, turning students into additional investigators.

The platform may participate in stabilization as well. Card `SCOTT-COP-006` records Roberts's public comment that Midjourney had interviewed him during the early period of Midjourney Magazine. The interview itself has not been recovered, so no influence claim is warranted. But the missing source points toward a significant mechanism: a platform can observe heterogeneous practitioner behavior, select exemplars, publish descriptions of those practitioners, and circulate a curated representation of the community back to the community.

A history of prompt engineering that archives only final successful prompts will therefore miss the practice. The stronger archive is longitudinal. It preserves failed prompts, remixes, who saw what, when a term first appeared, who copied it, which explanations accompanied it, how model versions changed outcomes, and when an informal procedure entered a syllabus, magazine, guide, or public talk.

The unit of evidence becomes a lineage of transformations rather than an isolated prompt-output pair.

7 Yelmo as a diagnostic object

Yelmo remains unusually useful even before the "Scott" identification is proven. The supplied notes use a repeatable yellow-Elmo-like character to test whether Midjourney can sustain a subject across scenes. The same notes identify the failure of multi-character composition: adding another character could alter the others. That historical practitioner problem can be placed beside later technical work without asserting influence.

Cards `YELMO-FORAGE-001` through `YELMO-FORAGE-004` split "character consistency" into distinct technical problems. Textual Inversion treats persistence as a learned subject representation. DisenBooth asks which features count as identity and which count as context. ConsiStory shows that consistent identity can emerge relationally across jointly generated images without a separately optimized subject representation. Isolation-attention work demonstrates a different failure: multiple subjects can contaminate one another through internal attention, so preserving two characters can require selectively blocking representational interaction.

The sequence changes the historical question. The 2022 practitioners did not merely lack a better prompt. They were encountering several not-yet-separated problems – identity representation, identity invariants, cross-frame consistency, and multi-subject interference – through one practical symptom: "it is not really the same dude each time."

That is precisely what makes practitioner archives valuable. A craft problem can precede the vocabulary that later decomposes it.

8 What would falsify the argument

The present argument would weaken substantially under several findings.

First, recovered course records could show that Roberts's 2023 pedagogy drew no meaningful practices, examples, vocabulary, or social methods from his 2022 generative-art activity. In that case the compressed chronology would remain biographical coincidence rather than a practice genealogy.

Second, controlled learning experiments could show that executable prompt sharing provides no distinctive advantage – or creates only shallow imitation with worse transfer – compared with non-executable demonstrations. "Executable repertoire" would then need to be qualified or abandoned.

Third, reconstruction of early Midjourney status could show that peer-recognized expertise tracked only prior professional art expertise rather than differential participation in the new platform. The "weeks-old old-timer" formulation would then overstate platform-native learning.

Fourth, move-level studies could find that generative outputs rarely introduce possibilities absent from human plans and function mostly as visualization of already selected ideas. The distributed-cognition claim would need to shrink accordingly.

Finally, the Yelmo identification remains a hard boundary. The present package retains the user-supplied Roberts biography next to a transcript that names a "Scott" described as an animator, but it does not promote that person to Scott Roberts without a primary bridge.

9 The next archive

The next evidentiary move is not another general theory of AI creativity. It is reconstruction.

Recover one complete collaborative Midjourney episode. Preserve all prompts, outputs, timestamps, participant messages, remix paths, and what each participant could see.

Recover one complete semester of the 2023 course. Preserve syllabus versions, assignments, tool introductions, student discoveries, instructor demonstrations, and project lineages.

Recover the missing Midjourney interview Roberts said had occurred.

Recover the full Surrealist Storytelling chapter and pre-2022 Roberts materials.

Those sources can turn executable apprenticeship from a plausible analytical wager into a historical account – or force the wager to change.

Evidence note

This paper is a derived interpretation, not evidence. Its claim-to-card-to-source trails are recorded in the accompanying `SOURCE_MAP`. External sources whose bodies were not stored are marked `LINK_ONLY`. The assembly prompt supplied for this compilation is preserved verbatim. Earlier conversation-native cards were reconstructed from the visible conversation rather than independently compared against a raw message export; the package records that limitation instead of claiming unverified byte-for-byte archival identity.

Selected references

- Edwin Hutchins. *Cognition in the Wild*. MIT Press, 1995.
- Jean Lave and Etienne Wenger. *Situated Learning: Legitimate Peripheral Participation*. Cambridge University Press, 1991.
- Jonas Oppenlaender. "A Taxonomy of Prompt Modifiers for Text-to-Image Generation." *Behaviour & Information Technology*, 2024.
- Etienne Wenger. *Communities of Practice: Learning, Meaning, and Identity*. Cambridge University Press, 1998.
- Scott Roberts. Public materials and posts registered in the accompanying source registry; external bodies remain `LINK_ONLY` unless locally preserved.