

THE MODEL COULD DRAW YELMO. IT COULD NOT KEEP HIM.

Scott Roberts, “the same dude each time,” and the minimum conditions of generative storytelling

Watson Hartsoe · Working Paper · 17 August 2026

Abstract

The original Yelmo project should not be remembered as a failed attempt to make a comic with Midjourney. Its proposal was more basic and more consequential: could an image generator sustain a protagonist long enough to tell a story? The surviving 2022 notes describe a small group imagining “a hero’s journey of Midjourney,” searching for “someone that’s kind of repeatable,” and settling on a deliberately wrong but recognizable figure: yellow Elmo, quickly compressed into Yelmo. The joke became a stress test. Across scenes, Yelmo’s eyes, frill, body, and relations drifted; adding other characters could cause them to alter one another. The problem was not simply visual inconsistency. The system could repeatedly generate descriptions of a type, but it did not carry forward a stable referent — *this same character* — from one image to the next.

This paper argues that Yelmo exposed a minimum ontology of generative storytelling before the field had good names for it. A story requires identity that survives transformation, not mere repetition; multiple actors whose boundaries survive interaction; and some mechanism that remembers what the story has already committed to. In the Yelmo experiments, much of that state appears to have lived outside the model, in the human group that named variants, rejected mutations, remembered prior versions, and treated a drifting family of images as one protagonist.

The model could draw Yelmo. The community was keeping Yelmo alive.

The argument becomes more interesting beside Scott Roberts’s documented practice. Years before generative AI, Roberts made work about entering the world of a cereal mascot, animating gamers’ idealized alter egos, crossing representation and physical space, and “experiment[ing] while making comics.” He later published a chapter titled “Surrealist Storytelling: Artistic Experimentation and Spontaneity through Improvisation,” publicly described Midjourney as a medium for an ongoing exquisite-corpse game in August 2022, and in 2023 taught an experimental AI-art course whose students he called research collaborators. None of this proves a direct lineage from those works to Yelmo. It does make a stronger question possible: was Yelmo not an AI novelty at all, but a new machine entering an older Roberts problem — how a character survives translation, improvisation, and transformation?

Keywords: generative AI; Midjourney; animation; character consistency; visual storytelling; improvisation; subject identity; Scott Roberts

1. The joke was already a theory of identity

The original proposal, as it survives in the research notes, is almost disarmingly small. A few people are trying to imagine “a hero’s journey of Midjourney.” They need a protagonist — “someone that’s kind of repeatable.” A friend named Scott proposes, or helps develop, an absurd test figure: yellow Elmo in a gritty farm-documentary register. The reason is partly

comic. Elmo is not yellow. “Yellow Elmo” rhymes its way into a name. Soon the group has Yelmo hiking a mountain, sitting at a fancy dinner, wandering into whatever scenario somebody can throw at him.

That choice is better than it looks.

Yelmo begins as a contradiction that remains recognizable. He is wrong before the experiment starts: Elmo recolored, category violated, identity bent. Yet he is still recognizable enough for the joke to land. The group has therefore chosen a subject that asks an unusually clean question: **how much can a figure change before it stops being the figure?**

That is not the usual benchmark for an image generator. Image generators are usually flattered by single-frame evaluation. Is this image beautiful? Does it match the prompt? Does the hand have five fingers? Is it photorealistic? Yelmo refuses that generosity. A gorgeous Yelmo on a mountain is not enough. The next image has to be *him again*.

The surviving notes state the failure with more precision than most technical vocabulary of the period: “it’s not really the same dude each time.” Sometimes the eyes change. Sometimes a red frill appears or disappears. The model seems to merge different yellow-Elmo-like possibilities. The image is good; the character is unstable.

That distinction is the beginning of the paper. **A model can succeed at depiction while failing at reference.** It can produce an image satisfying the description *yellow Elmo in an Italian restaurant* without possessing anything equivalent to the narrative commitment *the individual we already met on the mountain is now at this table*.

The difference sounds philosophical because it is. A story does not merely need recurring resemblance. It needs a way to say *this one again*.

2. The original proposal was not “make images.” It was “make a protagonist survive.”

The Joseph Campbell frame matters more than the notes give it credit for. A hero’s journey is not a slideshow of unrelated illustrations. Its basic wager is continuity through transformation. The hero leaves home, crosses thresholds, meets allies and enemies, is altered by ordeals, and returns changed. If the hero were visually frozen, nothing would have happened. If the hero became a different person every frame, there would be no one to whom the happening belonged.

So the original Yelmo proposal accidentally chose a much harder problem than “character consistency.” **It asked for identity-preserving transformation.**

That phrase matters. Consistency is often imagined as sameness: preserve the face, the shirt, the fur, the proportions. Narrative asks for something subtler. Change the pose. Change the lighting. Change the costume. Injure him. Age him. Frighten him. Put him underwater. Let him lose the coat that seemed visually definitive in the first image. Make him return from the underworld with something about him genuinely altered. And still: same dude.

This is why Yelmo is a better storytelling test than a static character sheet. The project forces a system to decide which differences are allowed to vary and which differences would destroy identity. Later subject-driven generation research would formalize versions of exactly this problem. Textual Inversion asked how a unique user-provided concept could be represented and recomposed into new scenes. DisenBooth explicitly separated identity-relevant information from pose, background, and other identity-irrelevant information. ConsiStory later pursued consistency across diverse prompts without requiring a permanently trained subject representation. These systems disagree about where persistence should live, but they share the problem Yelmo encountered experientially: a subject must survive variation without becoming trapped by its first depiction (Gal et al. 2022; Chen et al. 2023; Tewel et al. 2024).

The strange thing is that there is no neutral answer to what counts as identity-relevant. If every surviving Yelmo image has the same red frill, the frill begins to feel essential. If one image loses it and the group still says “that’s Yelmo,” the group has demoted it to costume. The reference set does not merely reveal identity; it helps legislate identity.

The original project therefore contains a question that should be put directly to Scott as an animator: **when you said “the same dude each time,” what did you actually mean by same?** Which properties did your eye forgive? Which mutations killed the character? Did you think in terms of model-sheet invariants — silhouette, facial spacing, proportion, costume, color — or was recognition looser and more performative? Was a bad drawing still Yelmo if the joke, pose, or attitude survived?

Those answers would matter more than another general interview about AI creativity. They would recover the human criterion against which the machine was being tested.

3. The community was the rig

Here is the braver claim.

In the workflow described by the notes, **Yelmo did persist. He just did not persist inside Midjourney.**

He persisted in the channel.

The people remembered what he had looked like. They named him. They compared new generations with prior generations. They laughed when the model produced an especially strange deviation. They rejected some mutations and accepted others. They proposed new situations for the same supposed individual. The continuity relation — *this image belongs to the Yelmo lineage* — was being maintained socially.

An animator will immediately see the inversion. A conventional character pipeline externalizes continuity into model sheets, turnarounds, palettes, rigs, reference drawings, naming conventions, asset libraries, and human review. Those devices make the character portable across shots and across hands. Yelmo had almost none of that inside the generative system. The group became the continuity apparatus.

The community was the rig.

This is not a metaphor for collaboration. It is a description of where state lived. The generator could be asked again for “yellow Elmo,” but the social group carried the history of what prior outputs had already established. Humans supplied the memory against which the next image became a continuation, mutation, or failure.

This makes the famous instability of early image generators newly interesting. Their lack of persistence did not simply prevent narrative. It recruited humans into maintaining the missing persistence themselves. The limitation helped produce the social game. Every broken Yelmo created an occasion to say: Is that him? Is this better? What if we try this? What if he goes here? The defect became a generator of conversation.

That may explain why the project could be simultaneously a failure as a production tool and a success as a social practice. The notes call Midjourney a “multiplayer imagination game.” Participants riff on one another, propose increasingly ridiculous scenarios, and watch the model make moves they would not have proposed themselves. The system is bad at maintaining a production asset and excellent at producing provocations.

The project’s own language catches the split: **game, not tool**. Not because nothing useful is happening, but because the labor of coherence still belongs to the people. A tool can be inserted into a pipeline because its outputs can be made answerable to prior decisions. A game can be enjoyable precisely because prior decisions remain vulnerable to surprise.

Yelmo sits at that threshold.

4. The second character reveals the missing world

The strongest technical clue in the surviving notes is not the changing frill. It is the failure that appears when another character enters.

“We couldn’t add other characters.” Yelmo would begin altering them; the characters contaminated one another.

That is a different problem from one-character persistence. It suggests that the model was not merely failing to remember Yelmo. It was failing to maintain **independent entity boundaries** inside a shared scene.

A story requires relations. Yelmo needs an ally, antagonist, lover, stranger, mentor, crowd. But a relation is possible only if the relata survive the relation. If introducing a second figure causes attributes to leak until the figures converge, then the system does not yet have two stable actors; it has one image field negotiating a compromise.

Later research on multi-subject generation gives this failure surprisingly literal names. IR-Diffusion identifies interference among subjects and proposes “Isolation Attention” to prevent one subject from referencing another in ways that produce subject fusion. StoryBooth describes “self-attention leakage” in which one character begins to look like another and frames the problem specifically as an obstacle to visual storytelling. These papers are not evidence that the Yelmo group anticipated their mechanisms, and there is no influence claim here. They are

valuable because they let us rename the 2022 practical complaint (He et al. 2024; Singh et al. 2025).

Yelmo was discovering that a story needs object boundaries.

A beautiful image can blur identities and remain beautiful. A story cannot. If Yelmo shakes hands with a pirate and comes away half-pirate, the failure is not cosmetic. The world has lost the ability to say who did what to whom.

This is where the original project becomes more ambitious than most discussions of prompting. The minimum generative story world requires at least three operations:

1. preserve an entity across transformations;
2. preserve several entities simultaneously;
3. allow relations among those entities without fusing their identities.

The Yelmo group found all three by trying to make a ridiculous Muppet go on an adventure.

5. Scott Roberts makes the story harder, not easier

There is still a provenance problem that should not be hidden because Scott may be the person who can solve it.

The Yelmo transcript identifies a “Scott,” describes him as an animator and the more professionally experienced artist in the group, and places him inside the project. The researcher supplying this archive identifies that Scott as DePaul animator and professor Scott Roberts. Public evidence independently places Roberts deep in Midjourney experimentation in the same period: on August 31, 2022, he described work made in Midjourney that included results from an “on-going exquisite corpse game”; a collaborator had publicly called another August exchange an “exquisite corpse collaboration” with Roberts through Midjourney. But I have not found a public Roberts source that says “Yelmo.” That bridge should remain visible until Roberts confirms it or an original record does.

The interesting part is that Roberts’s documented pre-AI work makes Yelmo more, not less, plausible as serious artistic research.

In 2002, Roberts made *Lucky*, a projection sculpture that he describes as allowing him to “enter the world of the cereal mascot,” a childhood dream. The same year, *My Favorite Character Is a Wizard* drew from interviews with role-playing gamers about their favorite characters; Roberts says the animation brings those “idealized alter egos to life” while revealing the mannerisms and idiosyncrasies of their creators. His artworks repeatedly cross representational systems: a cartoon cat projected into architecture, a low-resolution digital model texture-mapped with high-resolution photographs, objects conceived as Boolean intersections, drawings used to imitate the slow movement of a motion-picture camera. His comics page opens with the plain declaration: “I experiment while making comics” (Roberts, “Comics”).

Then there is the later chapter title: “Surrealist Storytelling: Artistic Experimentation and Spontaneity through Improvisation.” We do not yet have the chapter body in this archive, so its

argument cannot be smuggled in from the title. But the conjunction is difficult to ignore: storytelling, experiment, spontaneity, improvisation. Those are also the terms needed to describe the Yelmo practice at its best.

This does not prove Yelmo descended from *Lucky*, the Wizard animation, or a Surrealist method. Chronology cannot do that work. But it gives us a better historical question than “When did Scott start using AI?”

The better question is: **what old problems did Midjourney suddenly make cheap enough, fast enough, and social enough for Roberts to attack in a new way?**

The mascot matters. The alter ego matters. The representation crossing into another world matters. The willingness to improvise matters. If Yelmo is Roberts’s project, then the most interesting reading may be that generative AI did not create a new artistic concern. It gave an old concern a machine partner that could mutate a character faster than a human group could stabilize him.

6. The hero’s journey failed into a better research project

The notes suggest that the initial Joseph Campbell project did not harden into a finished narrative. It “morphed into just developing that character.” This is usually the point where a project history becomes apologetic: the story was never completed; the technology was not ready; the group got distracted.

That is the wrong reading.

The story proposal failed into the experiment it actually needed to become.

Before asking whether Midjourney could tell twelve stages of a myth, the group had to establish whether the same hero could survive stage one and stage two. Yelmo reduced “Can AI tell a story?” to a prior question: **Can there be a recurring somebody for the story to happen to?**

That is methodological progress. It takes a vague cultural question and turns it into a testable production problem.

The project then found the next prior question: can the hero encounter somebody else without identities bleeding together? And another: which visual features are constitutive of identity, and which can change? And another: if the machine is most enjoyable when it surprises the group, how much control can be added before the improvisational value disappears?

This last tension is especially important for Roberts because the notes preserve both sides. On one side is the animator’s demand for intentionality: enough control to tell a story, define a character, and carry work into the world. On the other side is the pleasure of “being led by MJ,” of not trying to control everything, of receiving a visual move no artist in the group would have proposed.

That is not a temporary contradiction to be solved by better software. It is a durable artistic problem.

The more perfectly the system obeys, the less it can surprise. The more radically it surprises, the less it can sustain a world.

Yelmo is valuable because it places those demands in the same body. Keep him. Change him. Surprise me. Do not lose him.

That is a better research program than “AI-assisted animation.”

7. Midjourney eventually named the requirement

There is a clean historical afterimage.

Midjourney’s own documentation describes a Character Reference as a way to recreate a specific character in multiple images and maintain consistency across scenes. The company announced Character References on March 12, 2024 (Midjourney 2024; Midjourney, “Character Reference”). In later versions, Omni Reference similarly lets users carry characters and objects into new creations. Midjourney’s documentation now presents as a product feature almost exactly the capability the Yelmo group was missing in 2022.

No causal claim follows. There is no evidence here that Yelmo influenced Midjourney’s roadmap. The important point is convergence. An early practitioner group independently formulated a limitation that later became legible enough to receive a dedicated interface and product vocabulary.

The original notes even contain the product argument before the product existed: without character consistency, the images remain “a beautiful garden of images” trapped as one-offs; with a definable character that can be placed in new situations, comics, animation, poster series, and sustained storytelling open up.

That is more than user frustration. It is a requirements document spoken in the vernacular of play.

The braver historical claim is therefore this: **Yelmo was an informal benchmark before the product had a benchmark for the thing that mattered.**

Not image quality. Not beauty. Not prompt eloquence.

Can the character come back?

8. The experiment should be rerun, not memorialized

If the original Yelmo materials survive, the project is unusually well suited to reconstruction.

The goal should not be to make a nostalgic gallery. It should be to rerun the same problem across the history of the medium.

Recover the original prompts, outputs, timestamps, remix chains, private-server messages, and the October 12, 2022 Office Hours material noted in the archive. Recover the earliest image the group agreed to call Yelmo. Recover the first moment a participant said some variant was *not*

him. Recover the failed multi-character images. Recover the abandoned Campbell sequence, if it ever existed as a list of stages or prompts.

Then build the benchmark the group was accidentally inventing:

Persistence. Put Yelmo through twelve sharply different scenes. Ask whether independent raters believe the same character survives.

Transformation. Deliberately alter pose, costume, affect, age, lighting, medium, and environment. Determine which transformations preserve identity and which destroy it.

Re-entry. Remove Yelmo for several scenes, then bring him back. Does the system recover the same individual or merely another member of the type?

Multiplicity. Add one, two, then five recurring characters. Measure when identities begin to contaminate one another.

Interaction. Make characters touch, exchange objects, occlude each other, change clothes, or share props. This is where mere visual separation becomes world coherence.

Improvisation. Run two conditions: one optimized for maximum identity control, another allowing the generator much more freedom. Compare not just fidelity but the number of genuinely useful surprises introduced by the machine.

This would turn Yelmo into something rare: a longitudinal artistic benchmark whose question survives radical changes in model architecture.

The point would not be to prove that 2026 tools are better than 2022 tools. Of course they are better at many of these operations. The point would be to preserve the original human criterion and watch the medium slowly become able to satisfy it.

9. Scott, this is the part only you can break

If the Scott in the Yelmo record is Scott Roberts, then the next step is not another literature review. It is to let you damage this argument.

The paper needs answers to questions only the original participants can supply:

What was the proposal before Yelmo existed? Was there actually a planned Campbell sequence, or is “hero’s journey” a retrospective shorthand for a looser desire to make a story?

Who invented Yelmo? Who first typed the yellow-Elmo prompt, who named him, and was the fact that Elmo “should” be red part of the conceptual attraction or simply a joke that worked?

What did repeatable mean to you as an animator? What were the non-negotiable features? Which variations did you accept as performance, costume, or bad drawing, and which made you say: no, that is another character?

Did you ever make an actual model sheet? Did the group try to formalize Yelmo outside the prompt — reference boards, prompt templates, seeds, image prompts, Photoshop corrections, naming rules, canonical images?

When did the hero's journey become the Yelmo laboratory? Was there a particular failure that killed the narrative plan, or did the social pleasure of putting him in random situations simply become the better project?

What happened when other characters entered? Did you understand the failure as prompt competition, model confusion, compositing failure, or something else? Do the images survive?

How much did you want to lose control? The notes contain both the animator's frustration with insufficient intentionality and the pleasure of letting Midjourney lead. Were those opposite camps in the group, or two demands you personally wanted the tool to satisfy at once?

Was exquisite corpse already your method? Your August 2022 posts document Midjourney exquisite-corpse play, and your later published chapter names Surrealist storytelling and improvisation. Did that tradition consciously shape Yelmo, or is the resemblance something this paper is imposing after the fact?

Did Yelmo enter the 2023 DePaul course? Your public description of that course calls the students research collaborators exploring bleeding-edge tools. Did Yelmo, character consistency, collaborative prompting, or the game/tool distinction become examples or assignments?

What do you still have? The most valuable source may not be an interview answer at all. It may be a folder of ugly images, a forgotten Discord export, a screenshot of a prompt chain, or a note naming the twelve stages.

If those materials contradict the paper, good. The point is not to turn Yelmo into a myth. It is to recover what the experiment actually discovered.

Conclusion: the character existed in the channel before he existed in the model

The original Yelmo project asked a question more durable than the technology that failed it.

Can a generative image system sustain a protagonist through change?

In 2022, the answer in this small experiment appears to have been: not by itself. Midjourney could produce remarkable individual Yelmo images. It could vary him endlessly. It could surprise the group. But it could not reliably honor the accumulating commitments that made one generated image the continuation of another. It could render a character description without consistently maintaining a character history.

The humans closed the gap. They remembered. They selected. They named. They laughed at mutations and carried accepted ones forward. They supplied the distinction between variation and replacement. They held the entity boundaries that broke when other characters entered. They supplied enough continuity for Yelmo to become socially real even while he remained computationally unstable.

That is why the original failure matters.

The model could draw Yelmo. It could not keep him.

And for a while, the people did.

If Roberts’s older work belongs in this genealogy, the final irony is even better. An artist who had already entered the world of a cereal mascot, animated other people’s idealized alter egos, crossed cartoons into physical space, experimented through comics, and written about improvisational storytelling encountered a machine that could summon characters almost instantly but could not promise that the character would still be there in the next scene.

The brave version of the history is not that generative AI suddenly empowered an animator.

It is that an animator and his friends used a stupid joke to discover what the machine still lacked before the field had agreed on the name of the problem.

Yelmo was the joke.

“The same dude each time” was the research question.

And the unfinished hero’s journey may be the most important artifact precisely because it never got past the condition a story needs before the adventure can begin: **someone has to be able to leave one image and arrive in the next.**

Source note

The core Yelmo evidence is a researcher-supplied 2022 Midjourney-community transcript/notes file. It describes a “Scott” who is an animator and participant in Yelmo but does not, in the surviving excerpt, provide a surname. The researcher identifies this person as Scott Roberts. Public sources independently document Roberts’s Midjourney practice in August 2022, including exquisite-corpse collaboration, but no public source located for this revision explicitly names Yelmo. The identification therefore remains a direct question for Roberts rather than a concealed assumption.

References

- Chen, Hong, Yipeng Zhang, Simin Wu, Xin Wang, Xuguang Duan, Yuwei Zhou, and Wenwu Zhu. 2023. “DisenBooth: Identity-Preserving Disentangled Tuning for Subject-Driven Text-to-Image Generation.” arXiv:2305.03374. <https://arxiv.org/abs/2305.03374>.
- DePaul University. “Scott Roberts.” Faculty profile. <https://www.depaul.edu/faculty/scott-roberts>.
- DePaul University Center for Teaching and Learning. 2023. “28th Annual Teaching and Learning Conference.” May 5, 2023. <https://resources.depaul.edu/teaching-commons/events/teaching-learning-conference/Pages/28th-Annual-Teaching-and-Learning.aspx>.
- Gal, Rinon, Yuval Alaluf, Yuval Atzmon, Or Patashnik, Amit H. Bermano, Gal Chechik, and Daniel Cohen-Or. 2022. “An Image is Worth One Word: Personalizing Text-to-Image Generation using Textual Inversion.” arXiv:2208.01618. <https://arxiv.org/abs/2208.01618>.
- He, Huiguo, Qiuyue Wang, Yuan Zhou, Yuxuan Cai, Hongyang Chao, Jian Yin, and Huan Yang. 2024. “Improving Multi-Subject Consistency in Open-Domain Image Generation with Isolation and Reposition Attention.” arXiv:2411.19261. <https://arxiv.org/abs/2411.19261>.
- Midjourney. 2024. “Introducing Character References.” March 12, 2024. <https://updates.midjourney.com/character-refs/>.
- Midjourney. “Character Reference.” Documentation. <https://docs.midjourney.com/hc/en-us/articles/32162917505293-Character-Reference>.
- Midjourney. “Legacy Features.” Documentation. <https://docs.midjourney.com/hc/en-us/articles/33329788681101-Legacy-Features>.

- Research archive. 2022. “Yelmo / Midjourney community interview and field notes.” Researcher-supplied transcript/notes; unpublished.
- Roberts, Scott. 2002. “Lucky.” Ubutopia. <https://www.ubutopia.com/lucky.htm>.
- Roberts, Scott. 2002. “My Favorite Character Is a Wizard.” Ubutopia. <https://ubutopia.com/wizard.htm>.
- Roberts, Scott. “Comics.” Ubutopia. <https://www.ubutopia.com/comics.html>.
- Roberts, Scott. 2022. Instagram post describing “an on-going exquisite corpse game” made with Midjourney. August 31, 2022. <https://www.instagram.com/p/Ch6jYdhOw0j/>.
- Roberts, Scott. “Surrealist Storytelling: Artistic Experimentation and Spontaneity through Improvisation.” In *The Rose Metal Press Field Guide to Graphic Literature*, edited by Kelcey Ervick and Tom Hart. Rose Metal Press. Table of contents available at <https://rosemetalpress.com/wp-content/uploads/2023/02/FGtoGraphicLiteratureTableofContents.pdf>.
- Roberts, Scott. 2023. “The Cyborg Muse: The Potential of AI Tools to Enhance and Accelerate Creativity.” LinkedIn / DePaul Center for Teaching and Learning. https://www.linkedin.com/posts/scott-roberts-3a63523_28th-annual-teaching-and-learning-conference-activity-7059776060971769856-V8fR.
- Singh, Jaskirat, Junshen Kevin Chen, Jonas Kohler, and Michael Cohen. 2025. “Storybooth: Training-free Multi-Subject Consistency for Improved Visual Storytelling.” arXiv:2504.05800. <https://arxiv.org/abs/2504.05800>.
- Tewel, Yoad, Omri Kaduri, Rinon Gal, Yoni Kasten, Lior Wolf, Gal Chechik, and Yuval Atzmon. 2024. “Training-Free Consistent Text-to-Image Generation.” arXiv:2402.03286. <https://arxiv.org/abs/2402.03286>.
- @onetoomanyoptions. 2022. Instagram post: “a plate of food” — “an exquisite corpse collaboration with @mrscottroberts through #MidJourney.” August 16, 2022. <https://www.instagram.com/p/ChWLLqvuOXe/>.